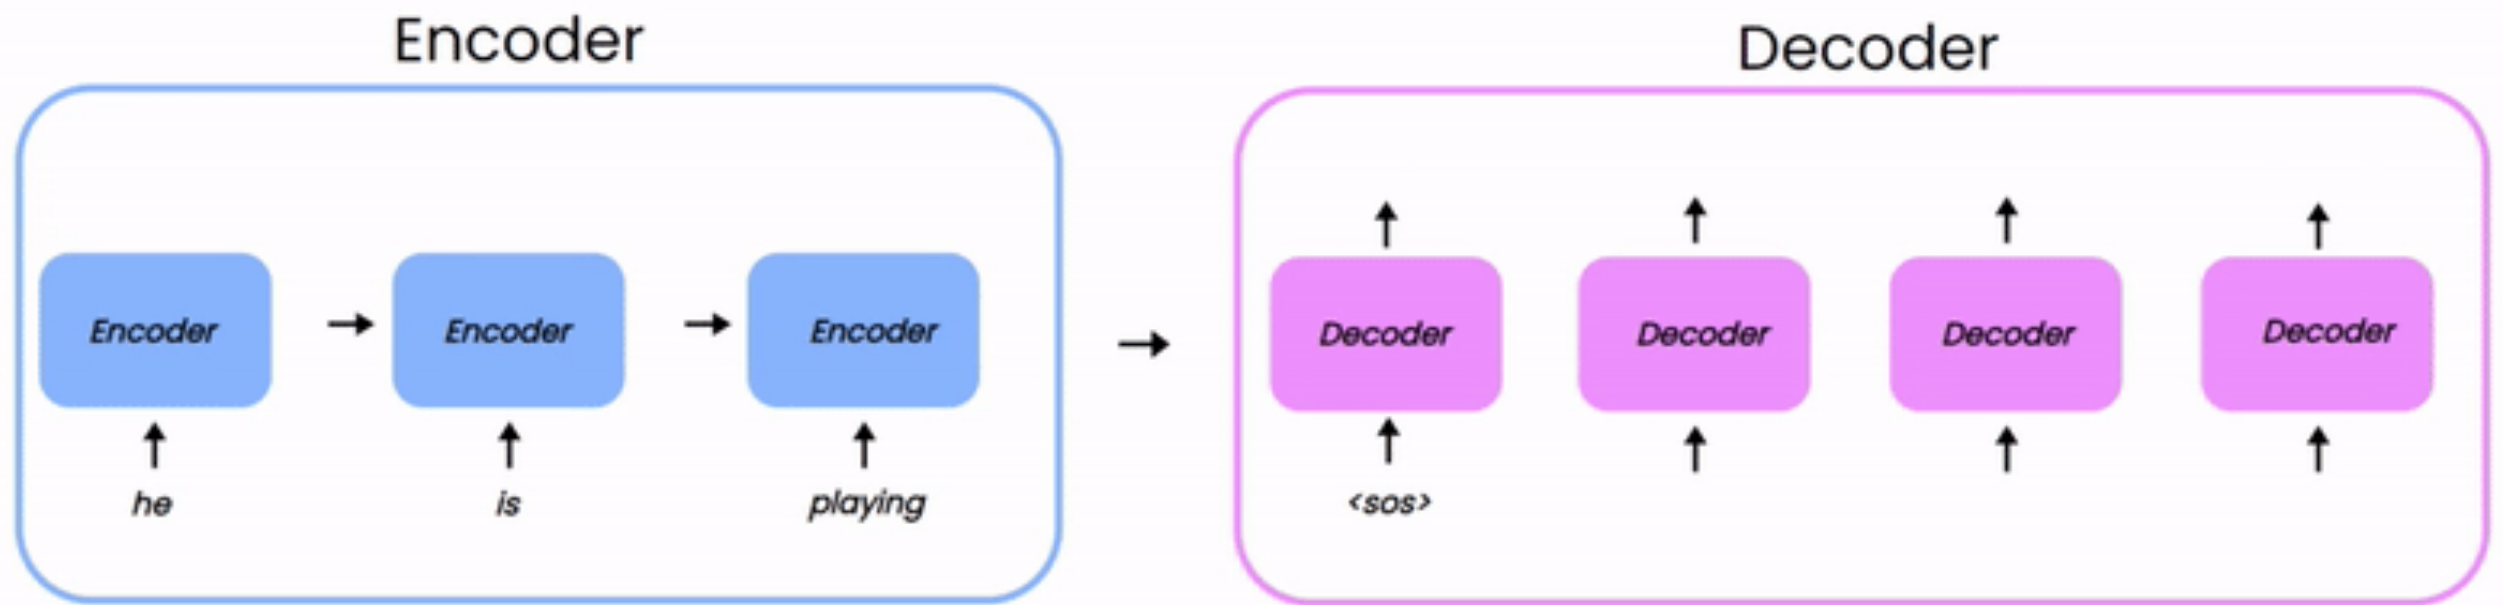


ENM 5310: Data-driven Modeling and Probabilistic Scientific Computing

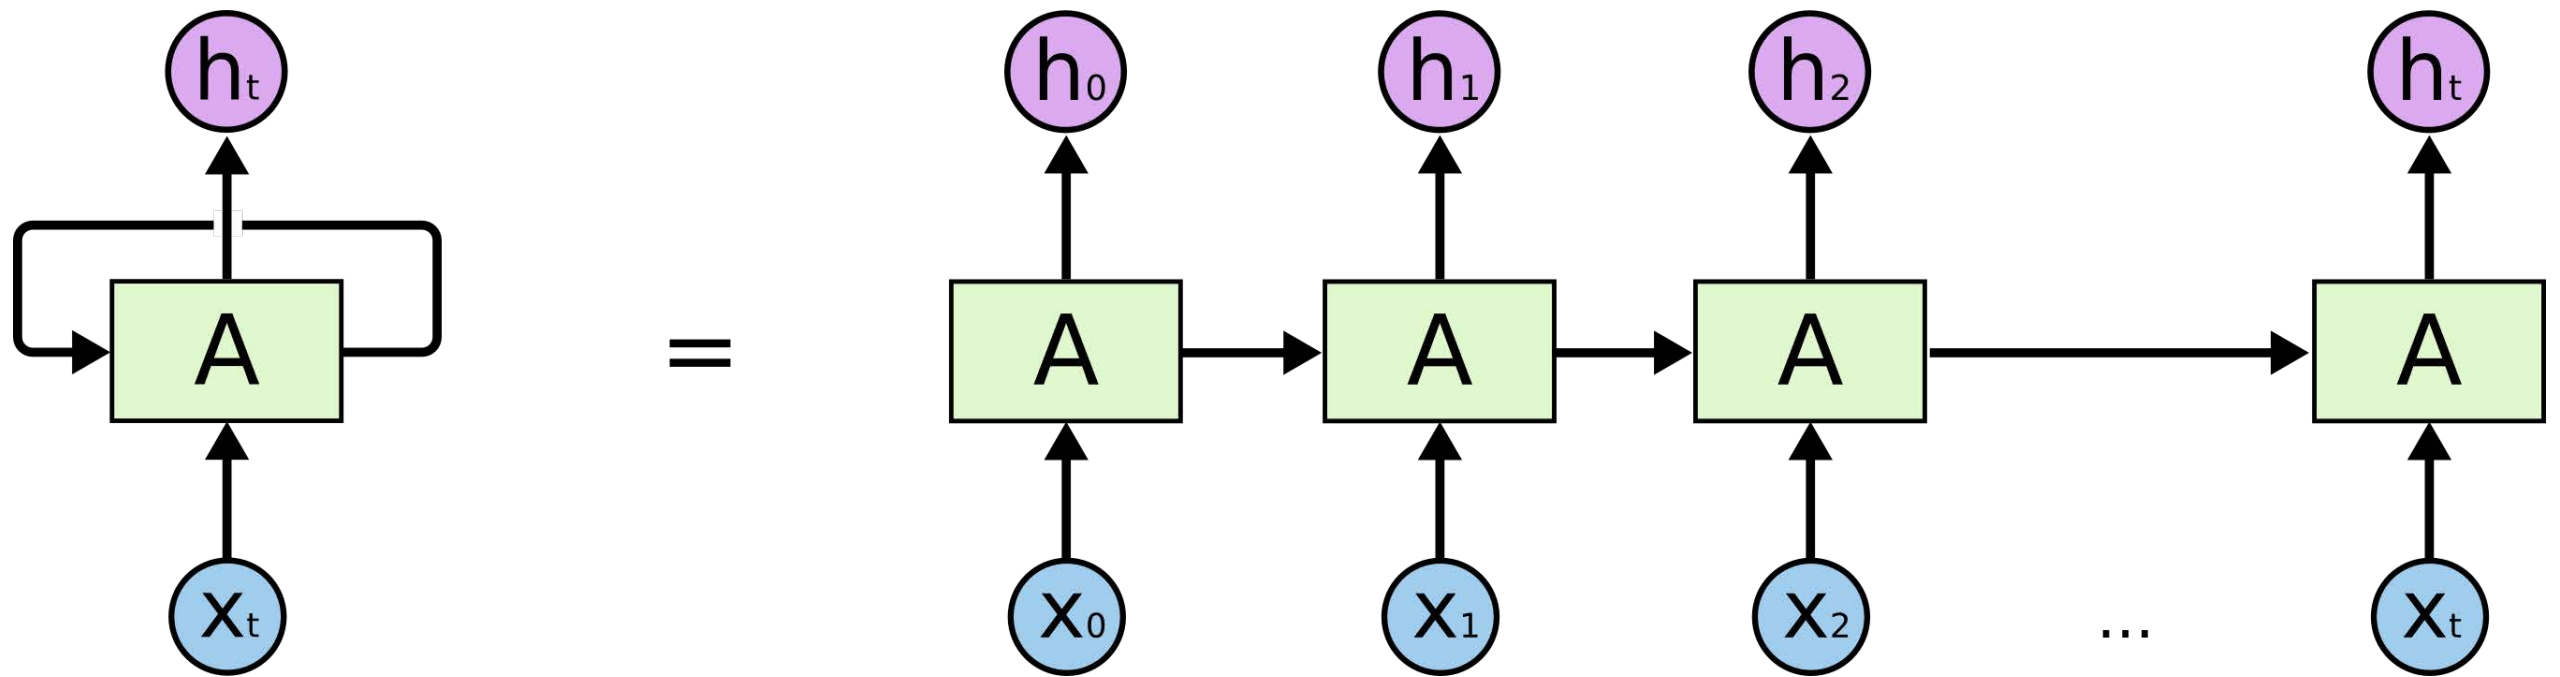
Lecture #15: Transformers




Modeling of sequence data



RNNs



An unrolled recurrent neural network.



RNN limitations

- Sequential processing is slow (can't parallelize)
- Long-range dependency issues (vanishing gradients)
- Limited context window in practice

*“The **cat** is on the **table**”*

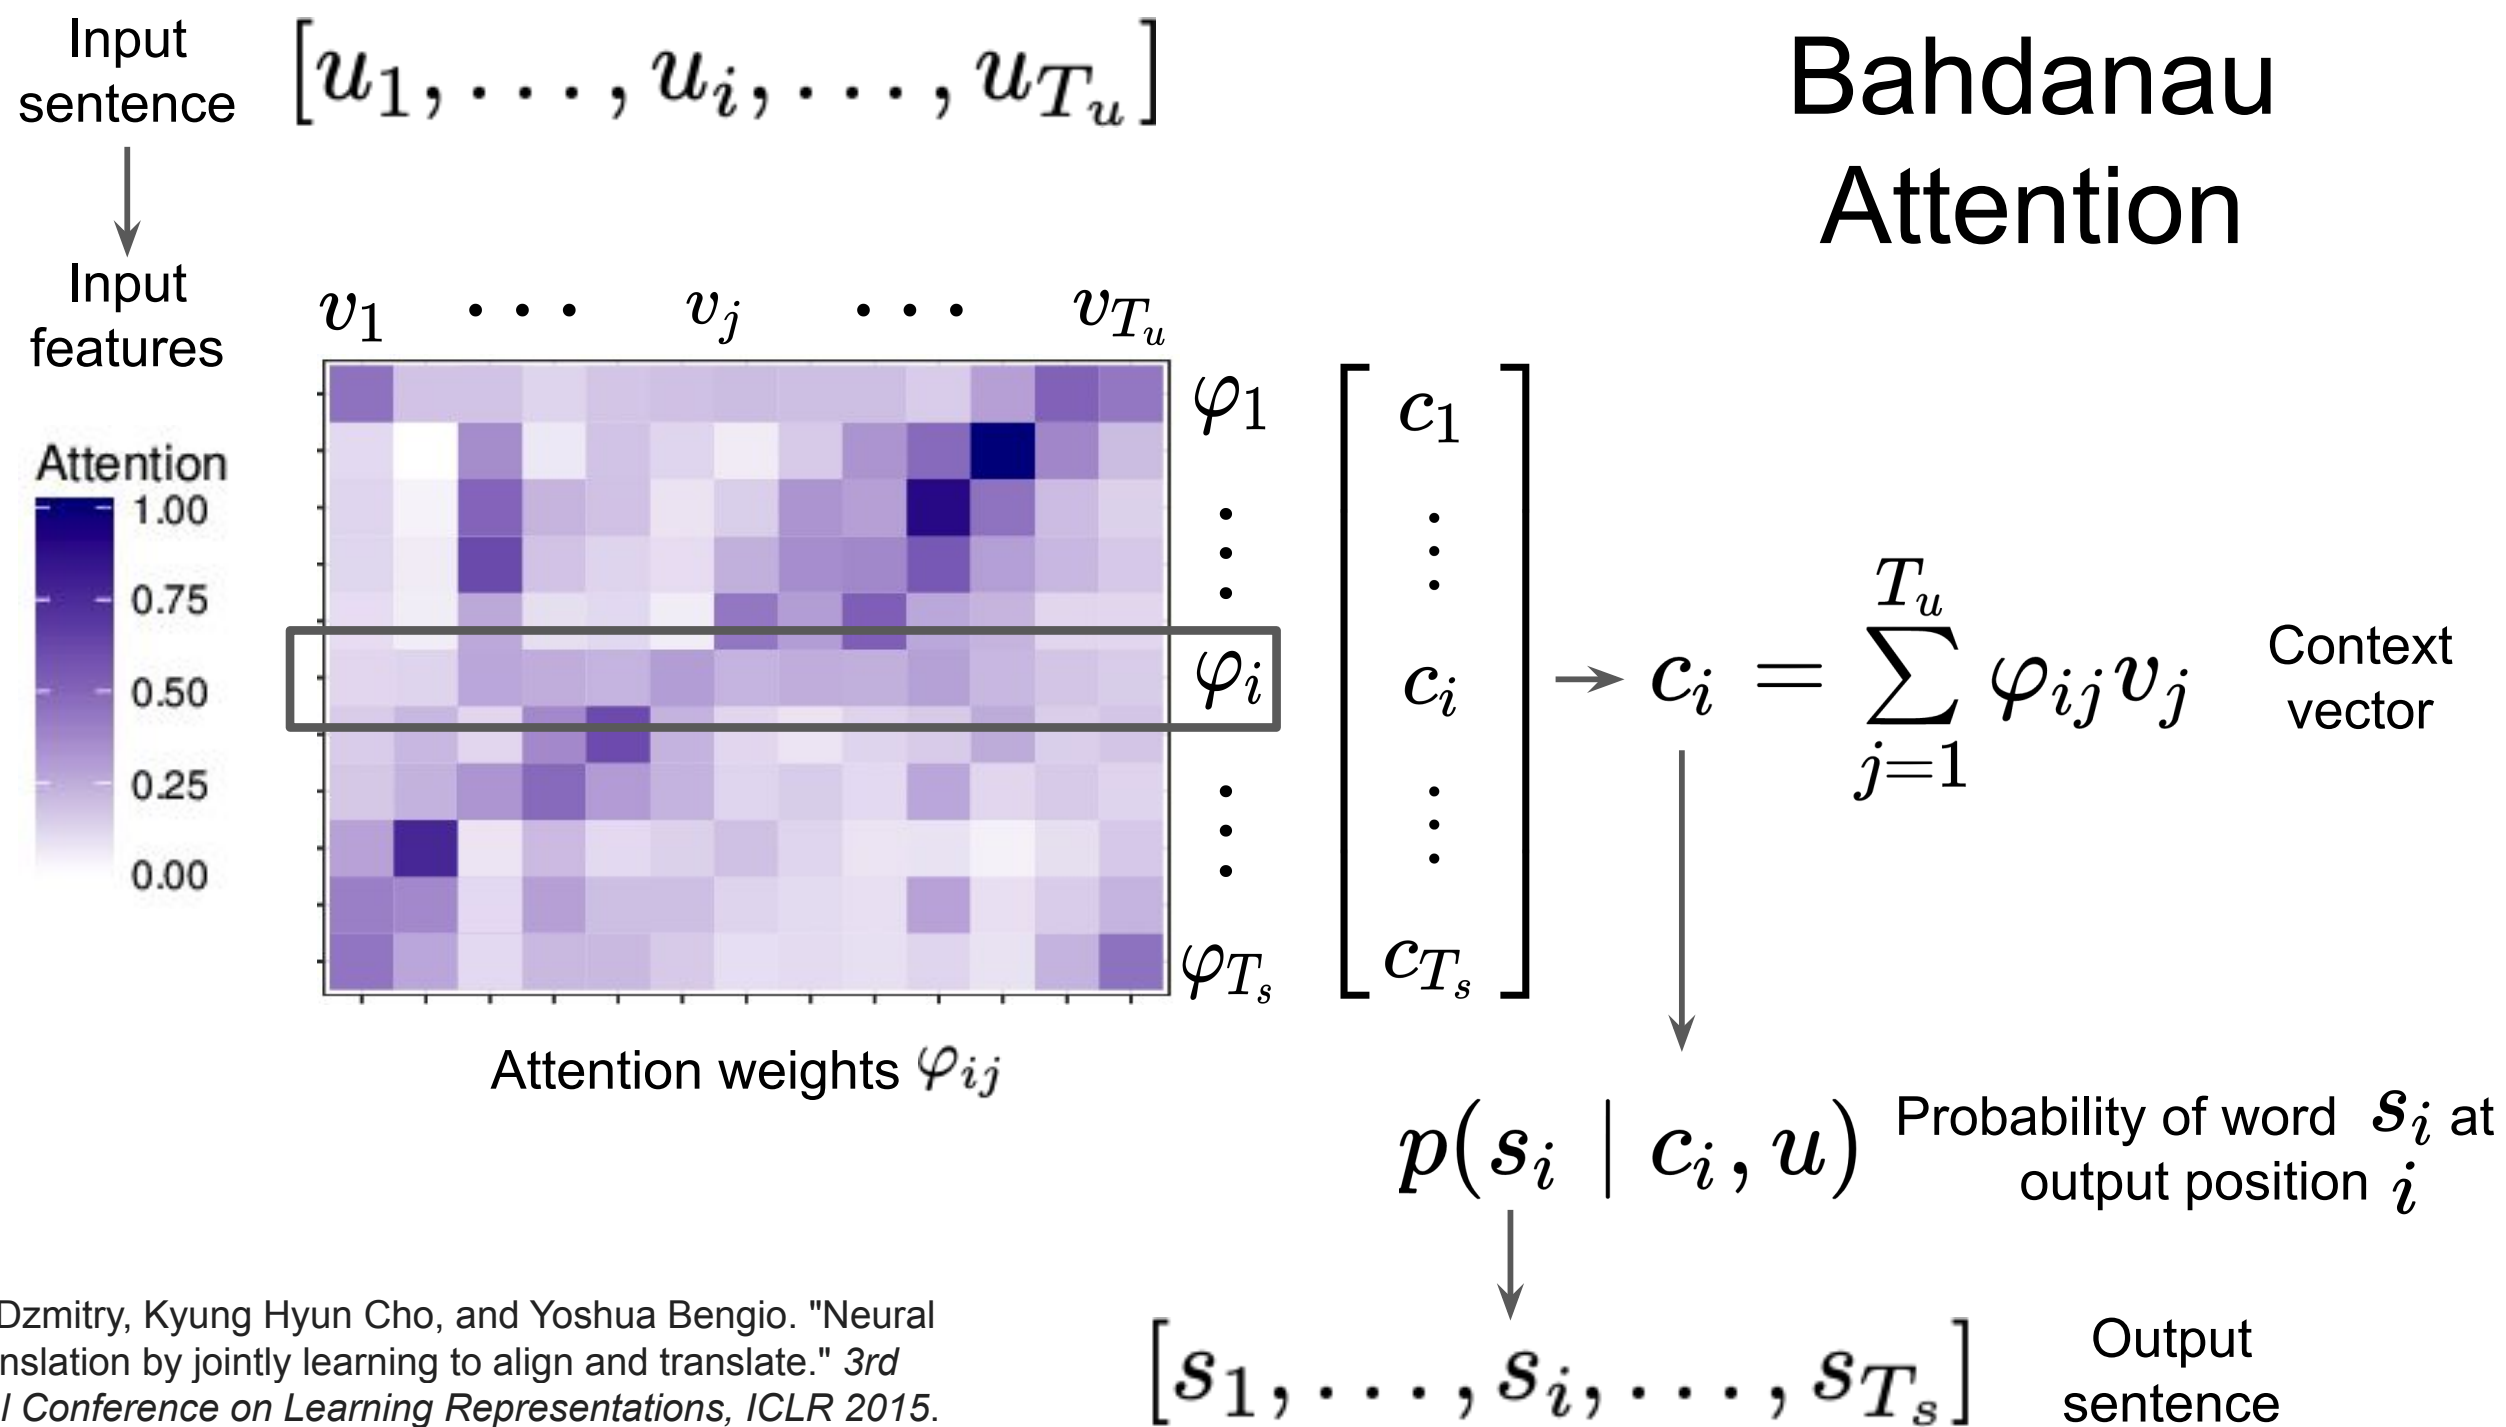
*“The **cat**, who belongs to my mother, is on the **table**”.*

E.g.: when translating we need to remember "cat" is the subject even after processing many words.

- A need for better models:
 - Language translation requiring full sentence context
 - Document summarization
 - Question answering
- Protein Structure Prediction (long amino acid sequences)
- Molecular Property Prediction (multi-atom interactions)
- Climate Science (long timescales)
- Astronomical Data Analysis (vast spatial and time scales)
-

Key Innovation: Attention

Bahdanau Attention



Bahdanau, Dzmitry, Kyung Hyun Cho, and Yoshua Bengio. "Neural machine translation by jointly learning to align and translate." *3rd International Conference on Learning Representations, ICLR 2015*. 2015.

From Sequential to Parallel Processing

- Compare how humans read vs how RNNs process text
 - Humans: Quick glances at relevant parts
 - RNNs: Must process word by word
- Attention is "looking" at all words simultaneously

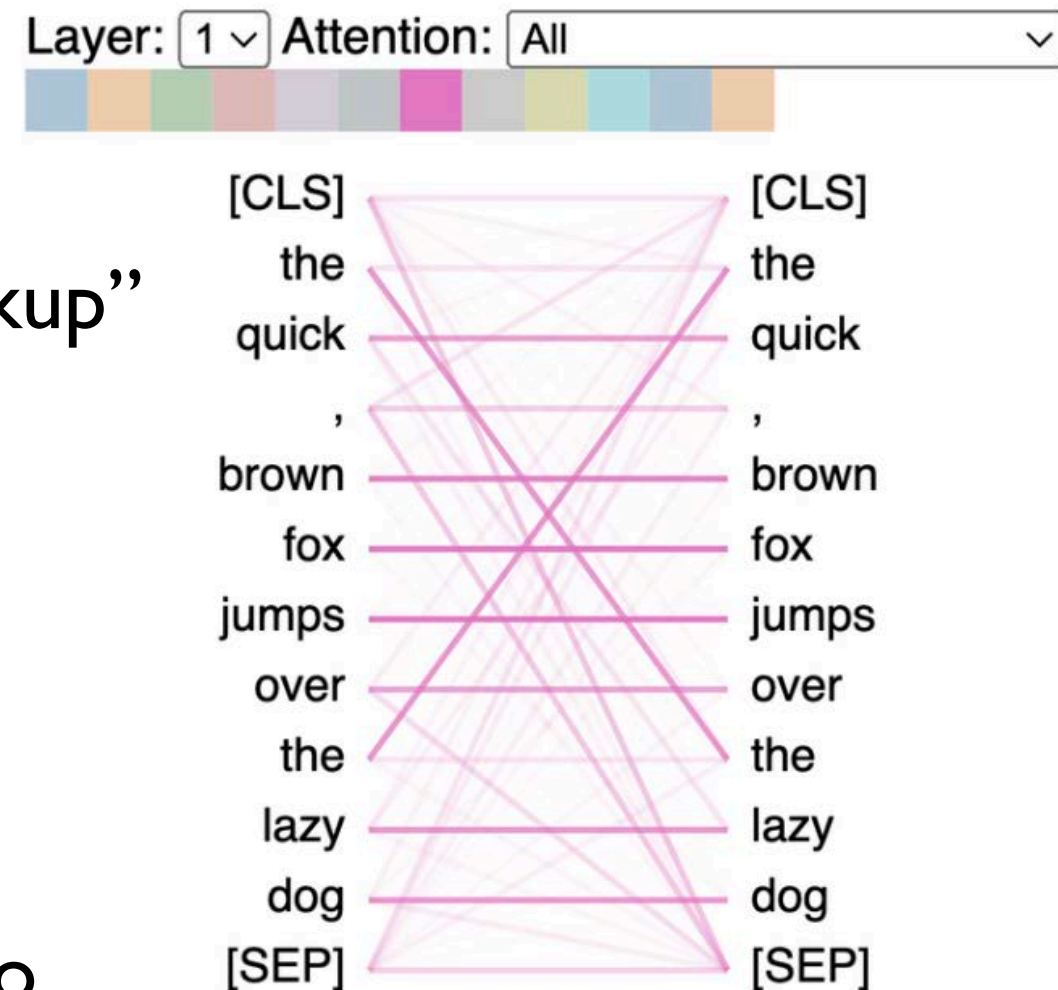
Key innovation: Attention

Think of attention as a "smart dictionary lookup"

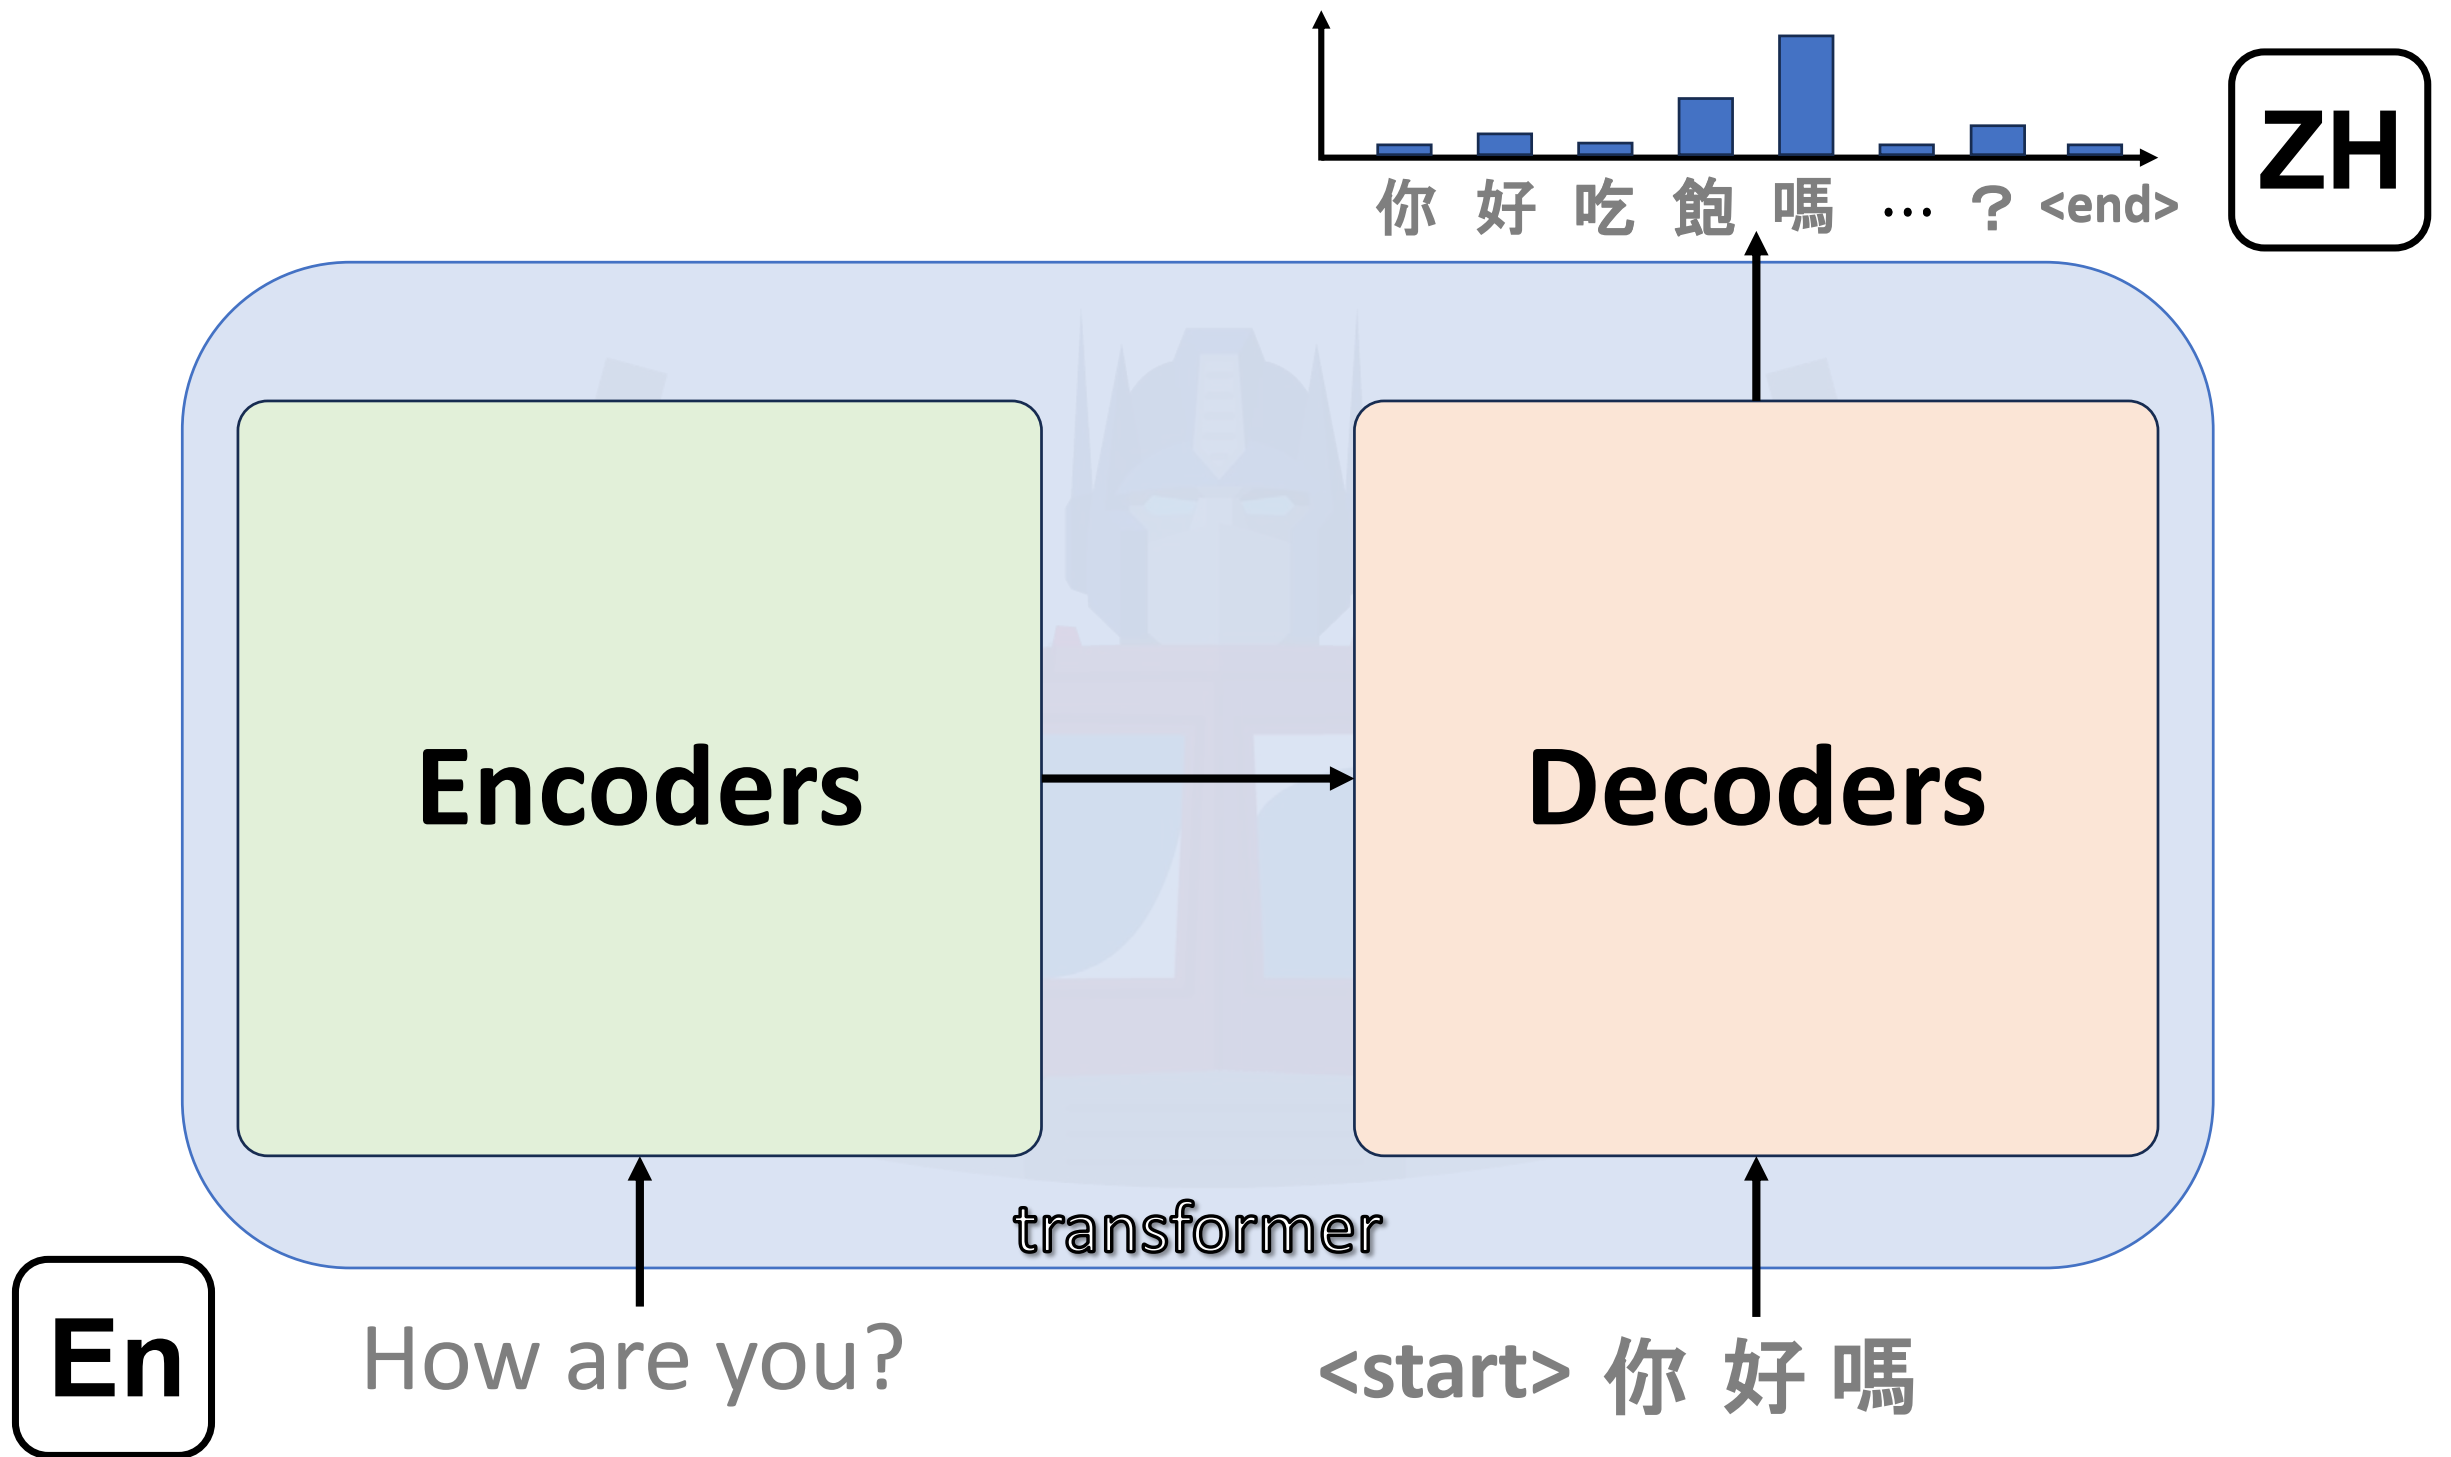
- Keys: What you're looking for
- Values: What you get back
- Queries: What you're asking for

Simple example: Translation attention

- Each word in target language might need to "look at" multiple source words



Transformers



Encoder-Decoder Auto-regressive Models

Key Building Blocks

- Tokenization
- Token embeddings
- Position embeddings
- Self-attention & Multi-head self-attention
- MLPs
- Cross-attention

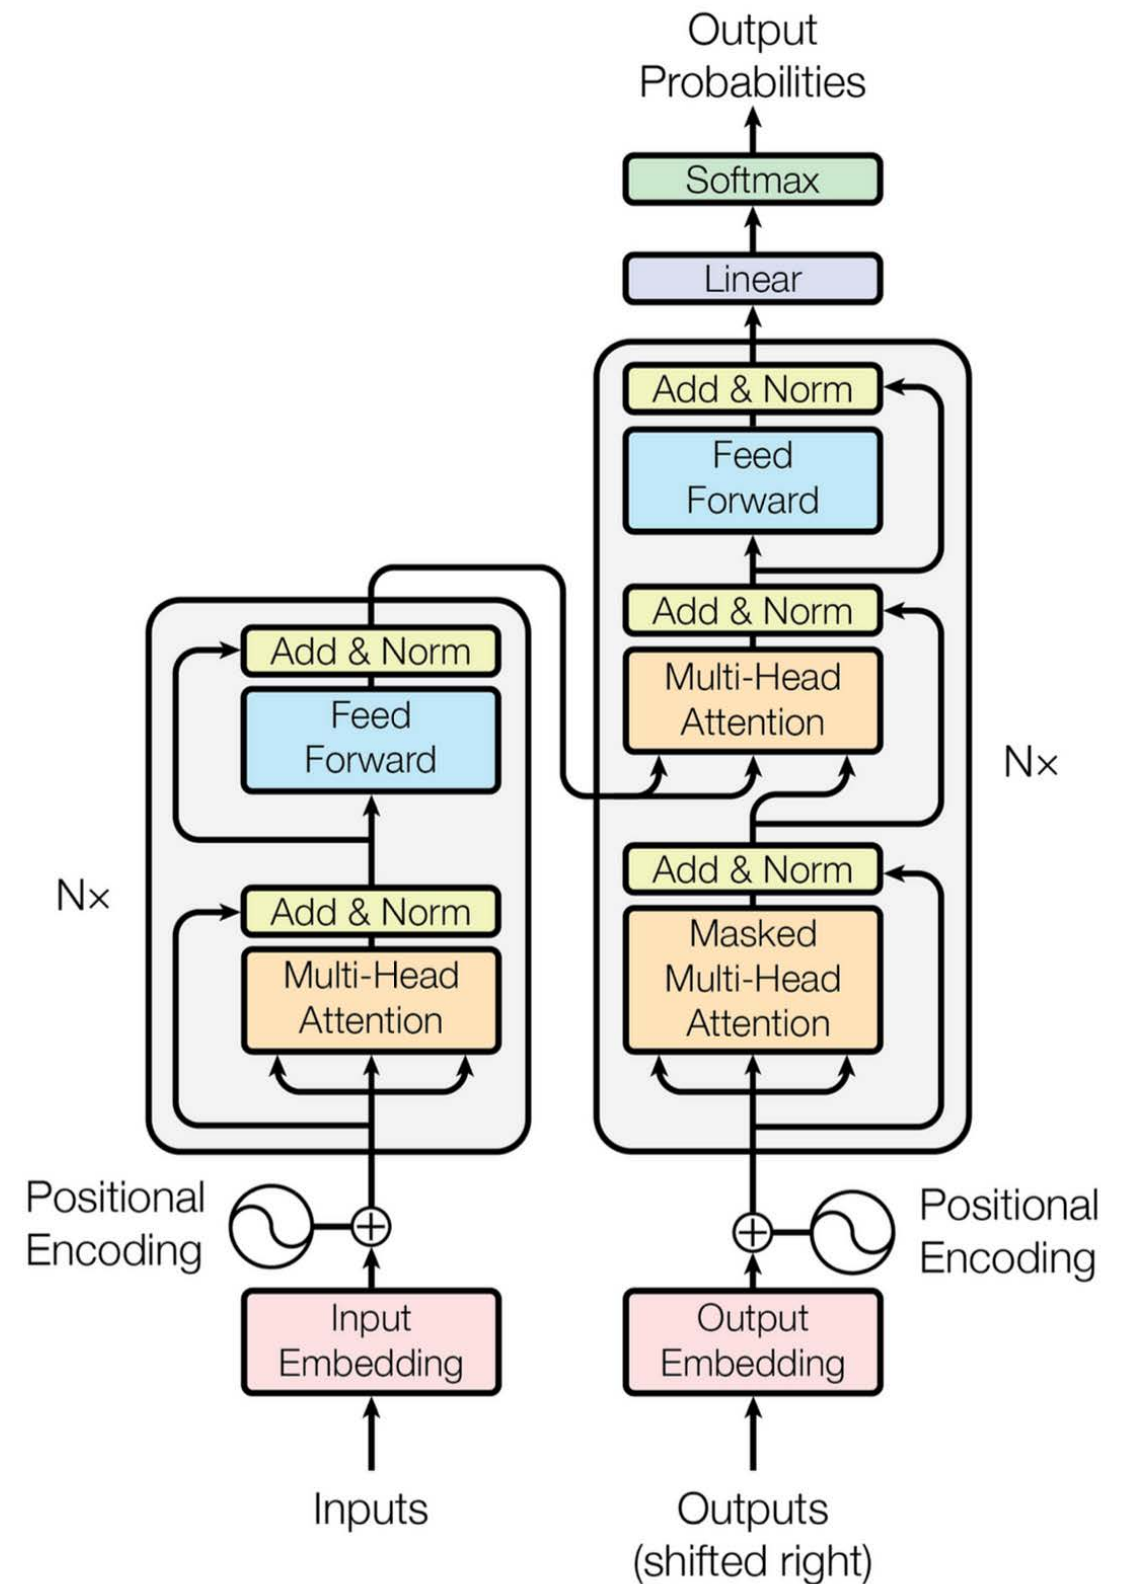




Figure 1: The Transformer - model architecture.

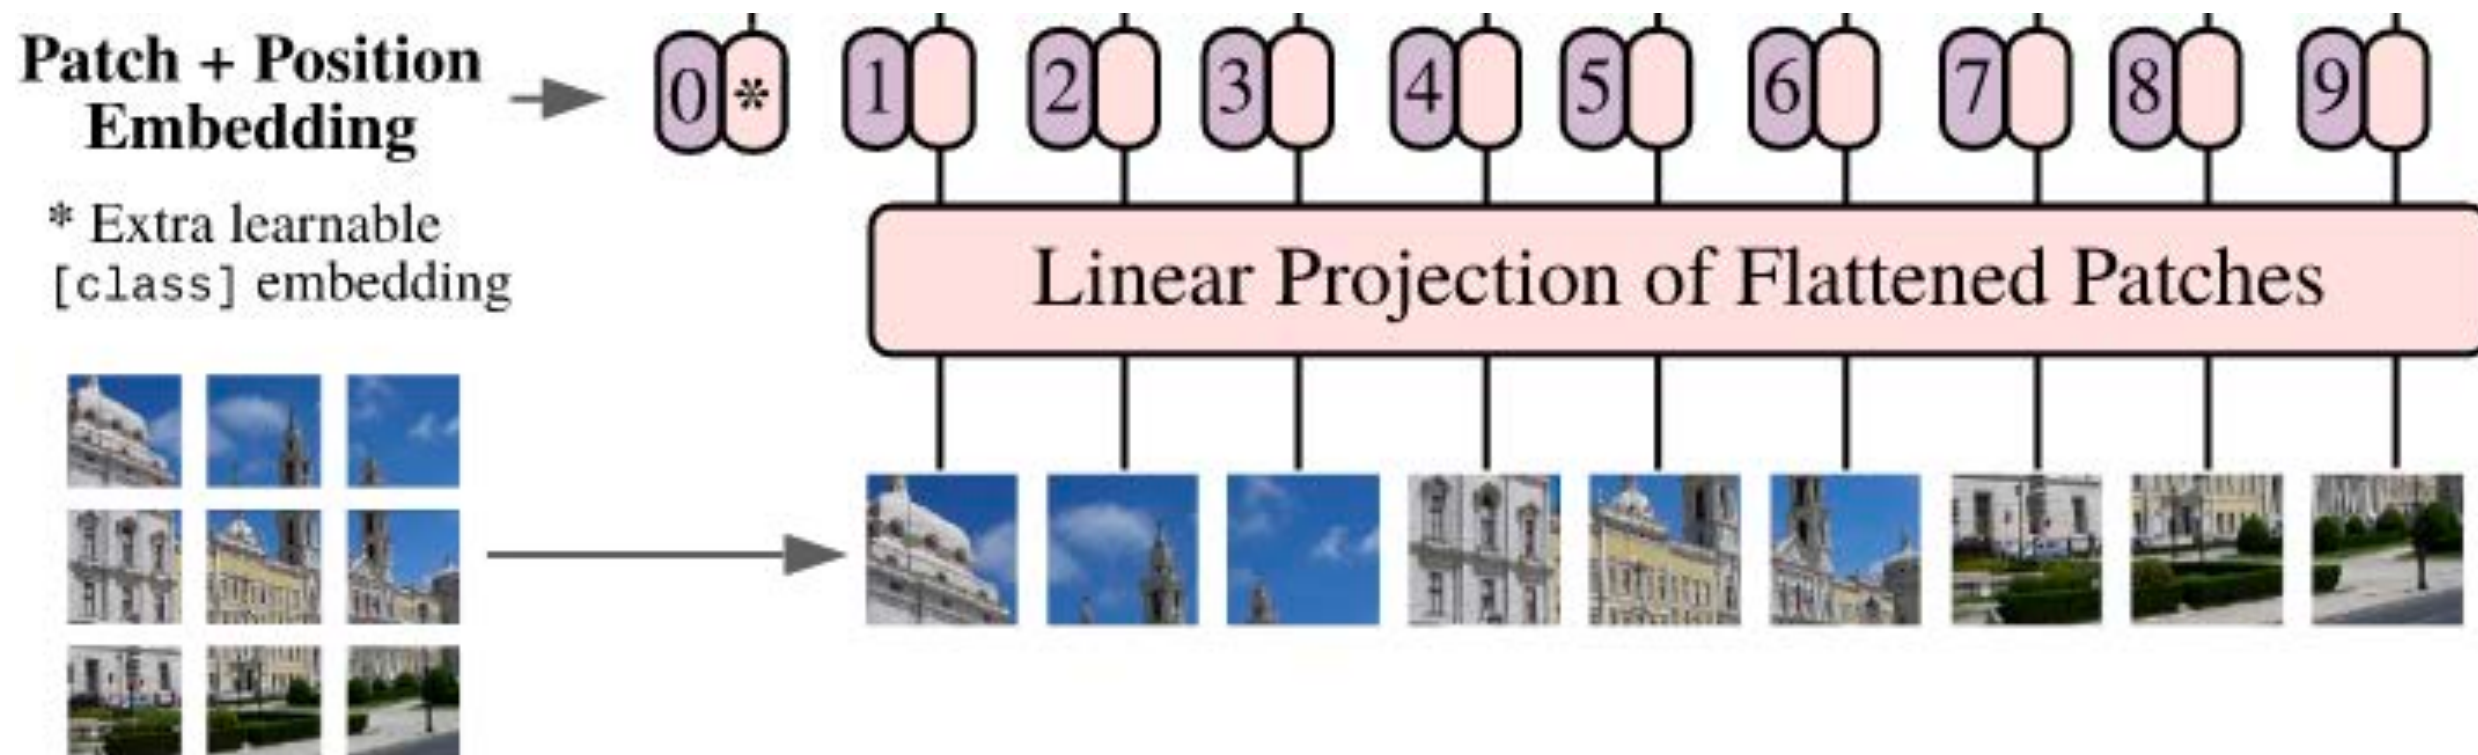
Tokenization in NLP




One-hot encoding




Tokenization in Vision




Tokenization in Point Clouds



Tokenization in NLP



One-hot encoding




Token Embeddings

cat	dog	bear	cow	indiv
$\begin{bmatrix} 1 \\ 0 \\ 0 \\ 0 \\ 0 \\ \vdots \\ 0 \end{bmatrix}$	$\begin{bmatrix} 0 \\ 1 \\ 0 \\ 0 \\ 0 \\ \vdots \\ 0 \end{bmatrix}$	$\begin{bmatrix} 0 \\ 0 \\ 1 \\ 0 \\ 0 \\ \vdots \\ 0 \end{bmatrix}$	$\begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \\ 0 \\ \vdots \\ 0 \end{bmatrix}$	$\begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ \vdots \\ 0 \end{bmatrix}$

Value 1 at 3687th entry

I bought an **apple** and an orange.

I bought an **apple** watch. .



Embedding Space

Apple

$$d \begin{bmatrix} 0.5 \\ 2.7 \\ 1.2 \\ \vdots \\ 0.2 \end{bmatrix}$$

Embedded token

$$= d \begin{bmatrix} \text{# tokens} \\ \vdots \end{bmatrix} \begin{bmatrix} 0 \\ 1 \\ 0 \\ 0 \\ 0 \\ \vdots \\ 0 \end{bmatrix}$$

Embedding Matrix


Token Embeddings

cat	dog	bear	cow	indiv
$\begin{bmatrix} 1 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ \vdots \\ 0 \end{bmatrix}$	$\begin{bmatrix} 0 \\ 1 \\ 0 \\ 0 \\ 0 \\ 0 \\ \vdots \\ 0 \end{bmatrix}$	$\begin{bmatrix} 0 \\ 0 \\ 1 \\ 0 \\ 0 \\ 0 \\ \vdots \\ 0 \end{bmatrix}$	$\begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \\ 0 \\ 0 \\ \vdots \\ 0 \end{bmatrix}$	$\begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ \vdots \\ 0 \end{bmatrix}$

Value 1 at 3687th entry

I bought an **apple** and an orange.

I bought an **apple** watch. .



Embedding Space

Apple

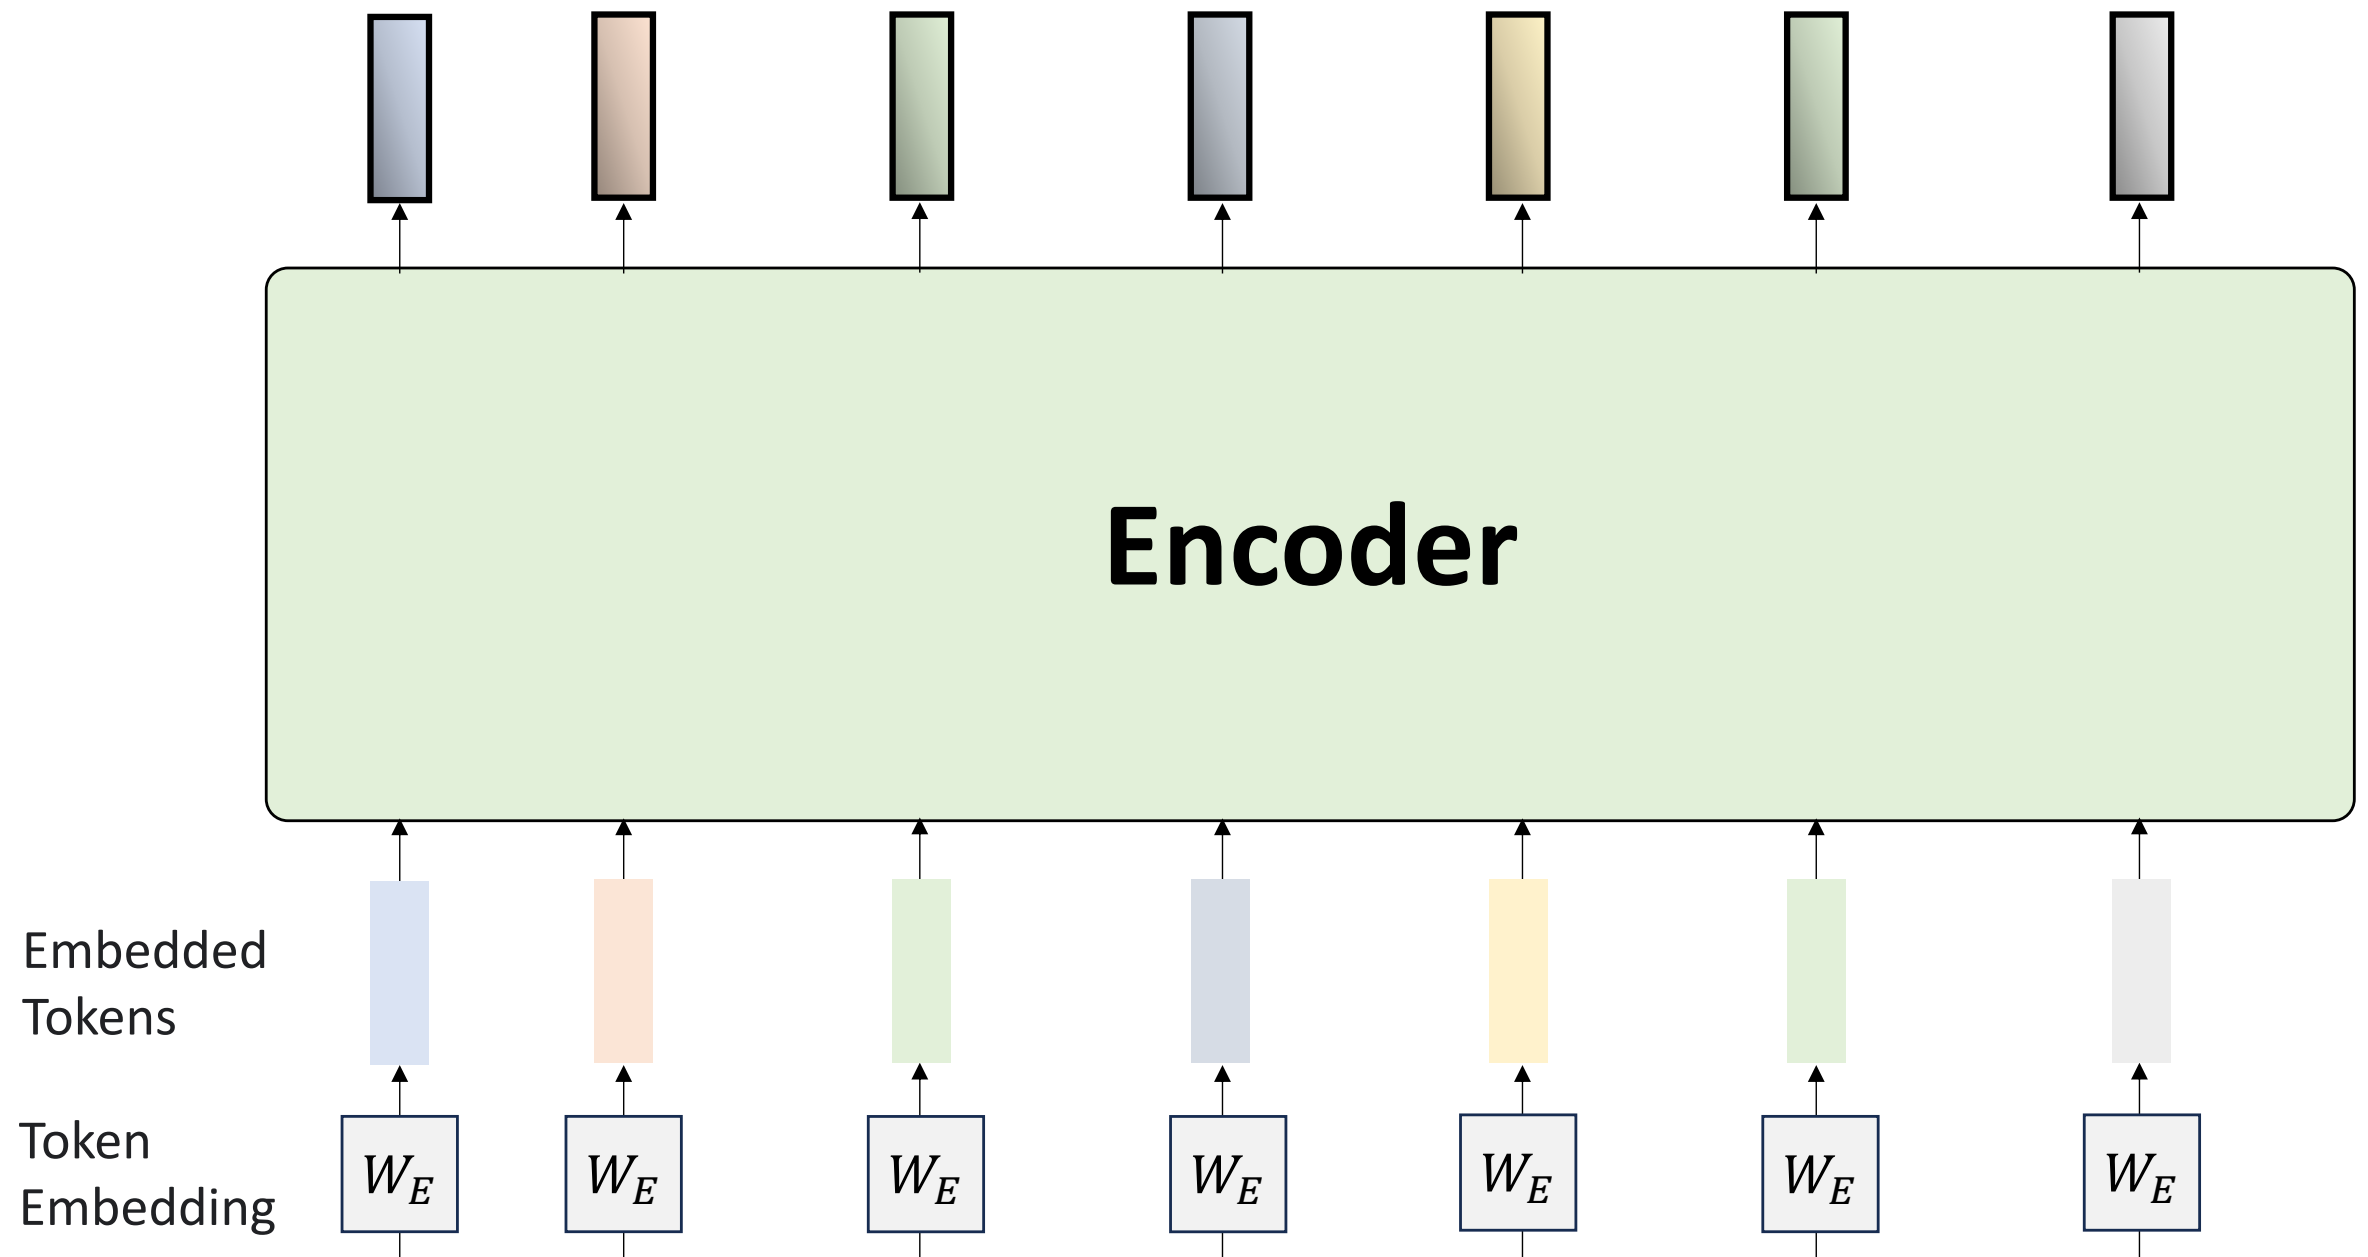
$$d \begin{bmatrix} 0.5 \\ 2.7 \\ 1.2 \\ \vdots \\ 0.2 \end{bmatrix}$$

Embedded token

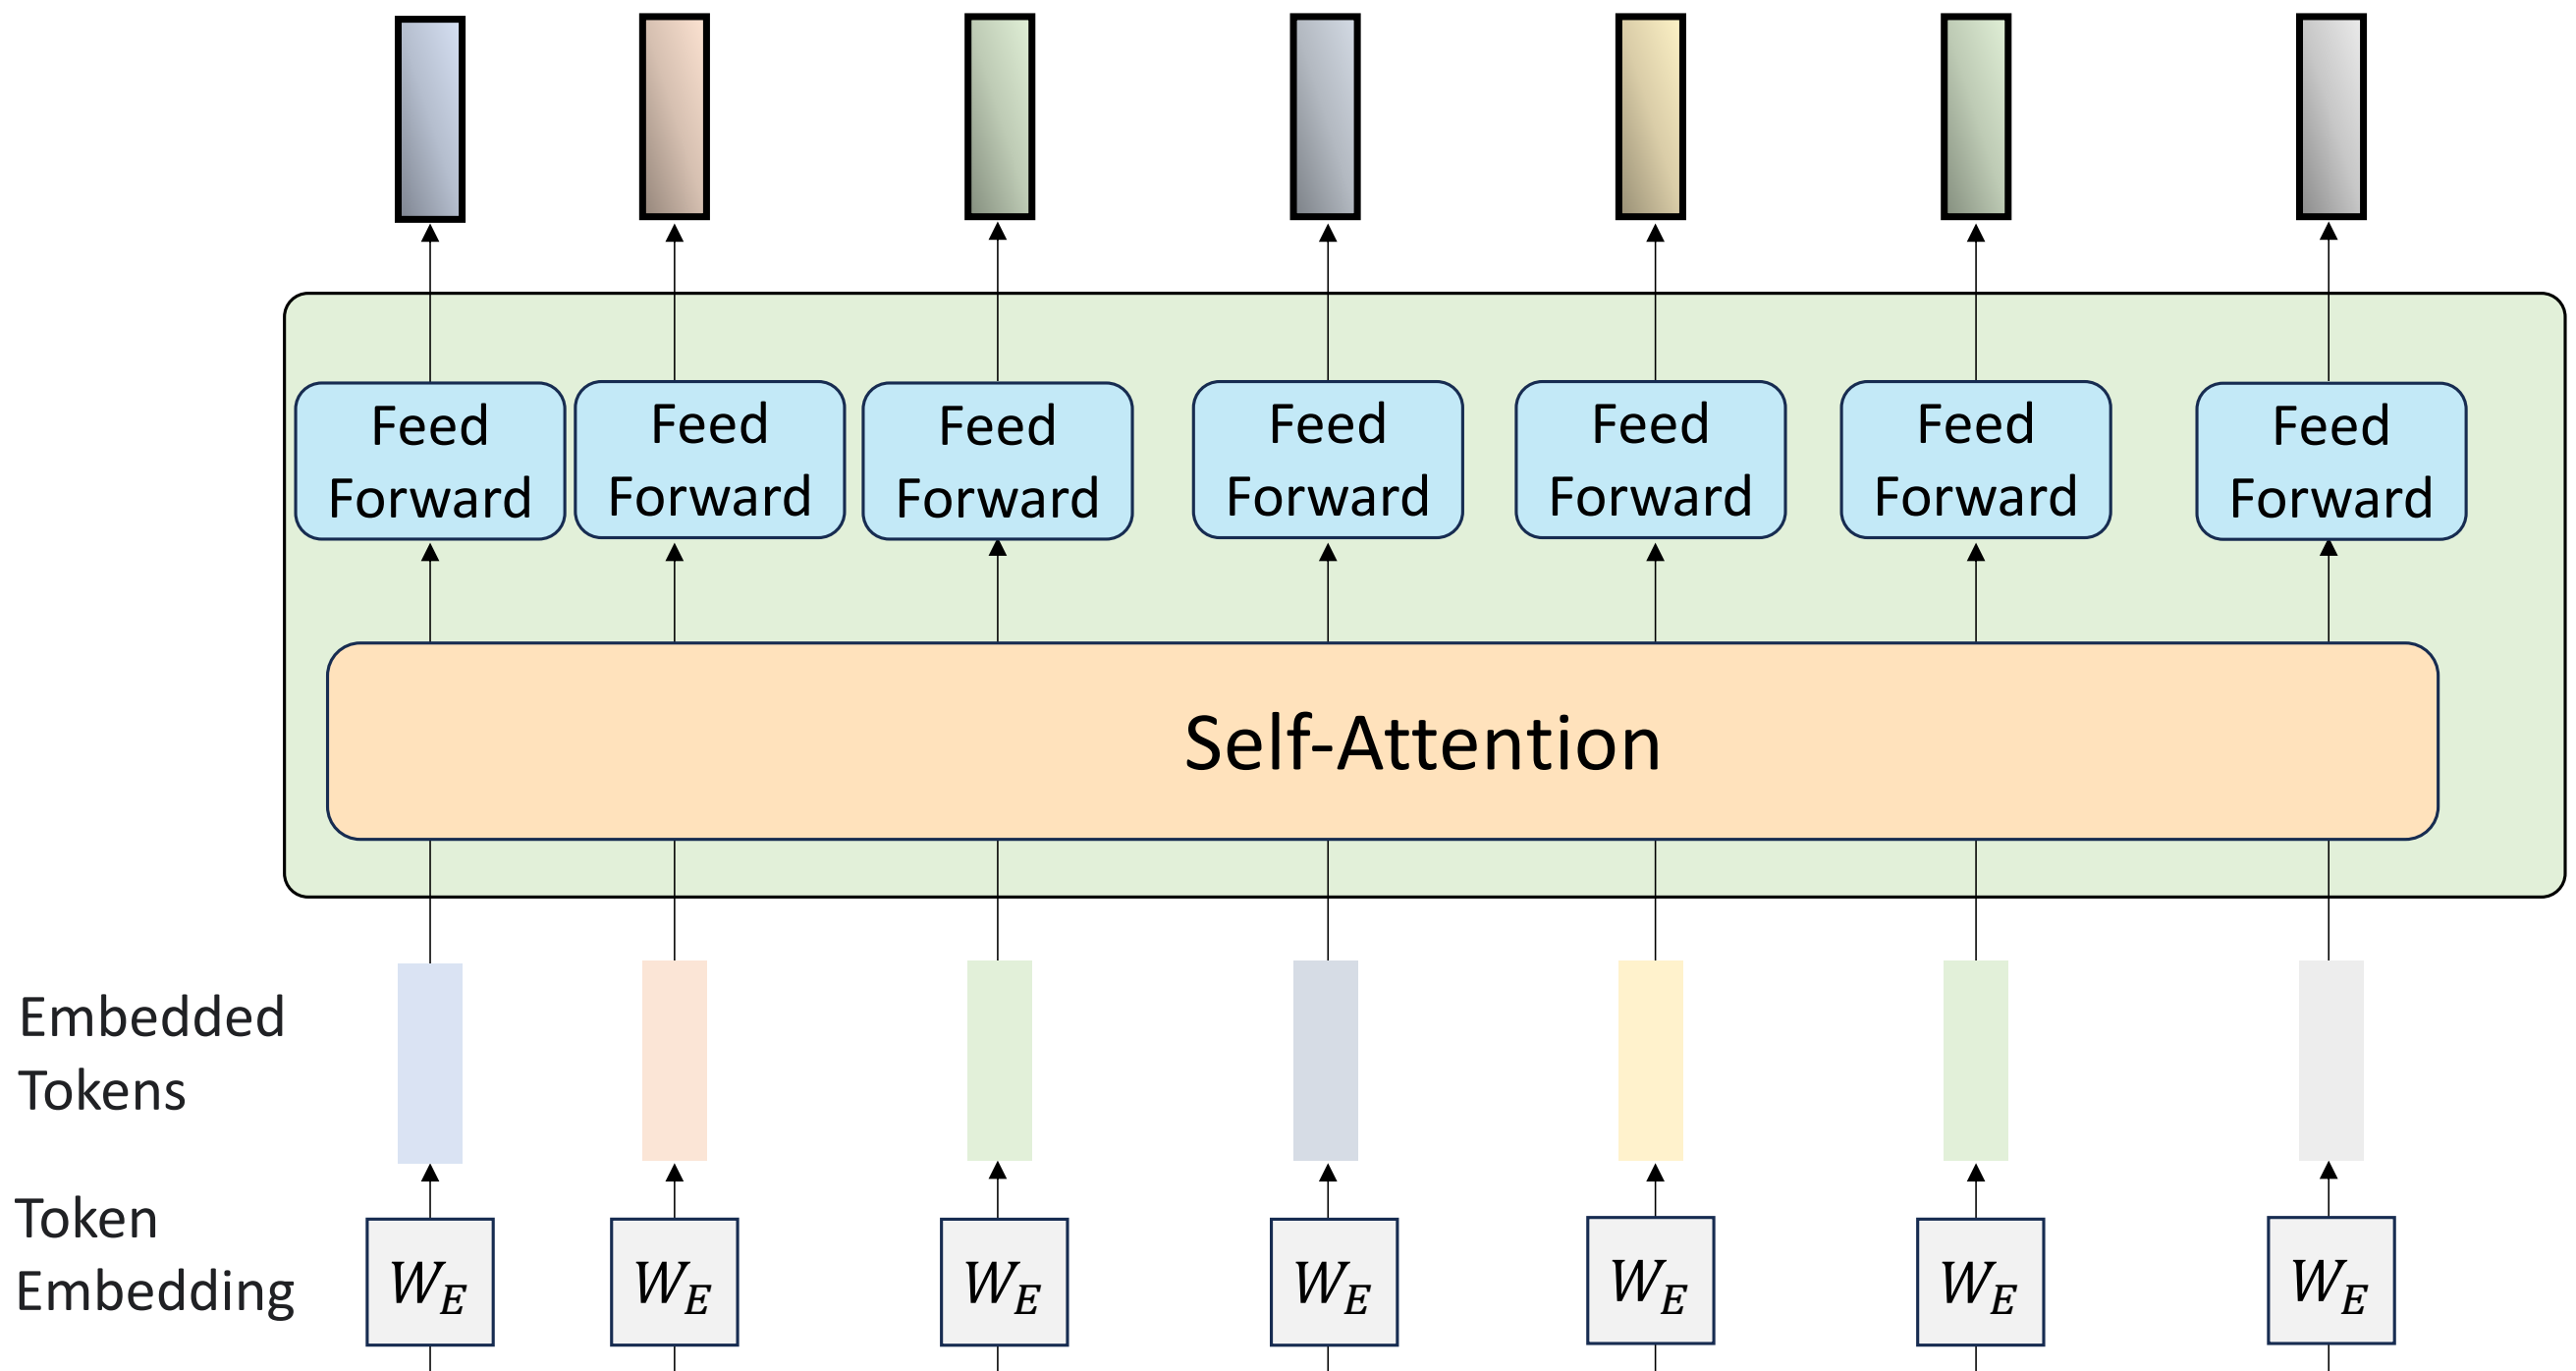
$$= d \begin{bmatrix} \text{# tokens} \\ \vdots \end{bmatrix} \begin{bmatrix} 0 \\ 1 \\ 0 \\ 0 \\ 0 \\ 0 \\ \vdots \\ 0 \end{bmatrix}$$

Embedding Matrix

Encoder



Encoder



Self-Attention

Updated feature $\mathbf{x}'_4 = \alpha'_{4,1} \mathbf{x}_1 + \alpha'_{4,2} \mathbf{x}_2 + \alpha'_{4,3} \mathbf{x}_3 + \alpha'_{4,4} \mathbf{x}_4 + \alpha'_{4,5} \mathbf{x}_5$

Attention Scores

$\alpha'_{4,1}$	$\alpha'_{4,2}$	$\alpha'_{4,3}$	$\alpha'_{4,4}$	$\alpha'_{4,5}$
0.082	0.0495	0.0199	0.6034	0.2452


Softmax

$$\alpha'_{4,i} = \frac{\exp(\alpha_{4,i})}{\sum_j \exp(\alpha_{4,j})}$$

Token similarity

$\alpha_{4,1} = \mathbf{x}_4^\top \mathbf{x}_1$	$\alpha_{4,2} = \mathbf{x}_4^\top \mathbf{x}_2$	$\alpha_{4,3} = \mathbf{x}_4^\top \mathbf{x}_3$	$\alpha_{4,4} = \mathbf{x}_4^\top \mathbf{x}_4$	$\alpha_{4,5} = \mathbf{x}_4^\top \mathbf{x}_5$
---	---	---	---	---

Embedded
Tokens



Self-Attention

$$\text{Updated feature } \mathbf{x}'_4 = W^O \left(\alpha'_{4,1} v_1 + \alpha'_{4,2} v_2 + \alpha'_{4,3} v_3 + \alpha'_{4,4} v_4 + \alpha'_{4,5} v_5 \right)$$

$$= \sum_i \alpha'_{4,i} \left(W^O W^V \right) \mathbf{x}_i$$

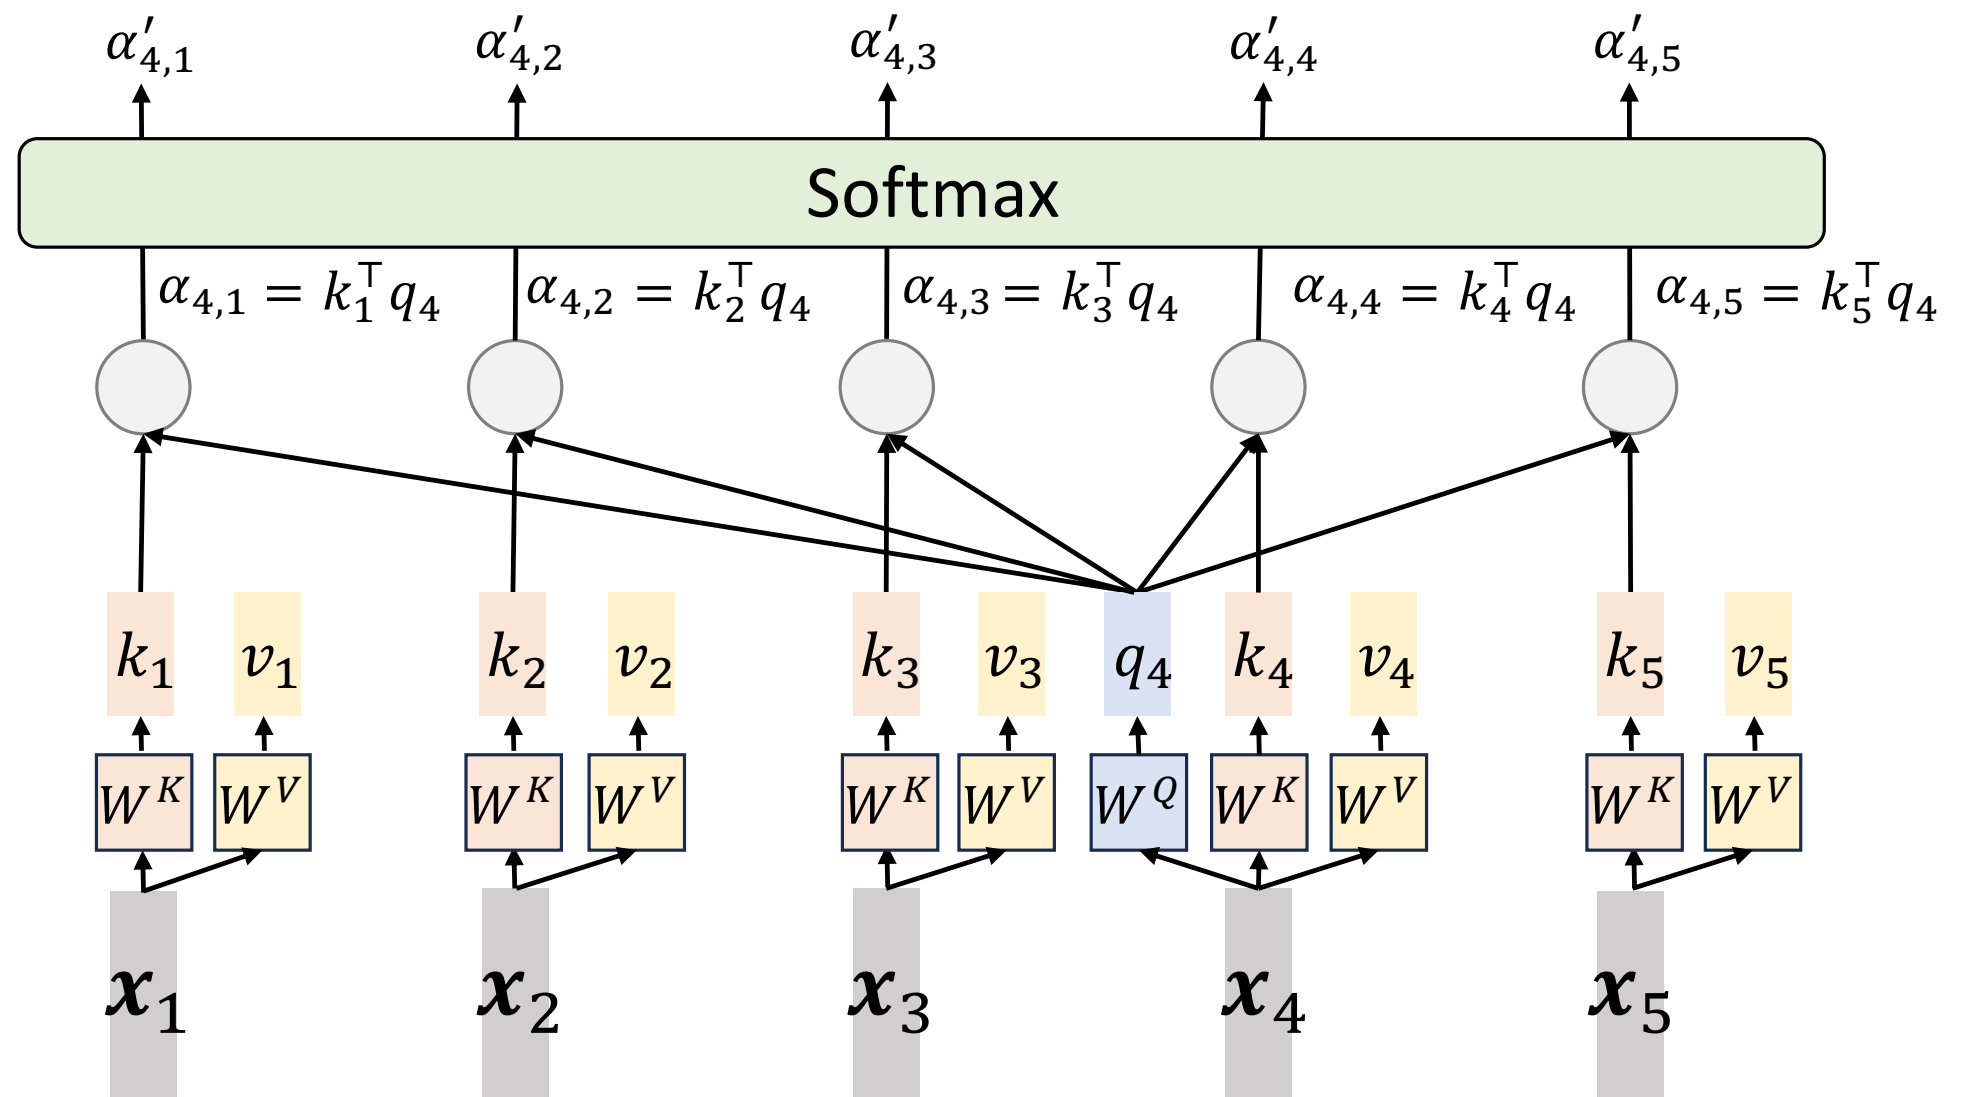
$$W^O \in R^{d \times d_v}$$

$$W^Q \in R^{d_k \times d}$$

$$W^K \in R^{d_k \times d}$$

$$W^V \in R^{d_v \times d}$$

Embedded
Tokens



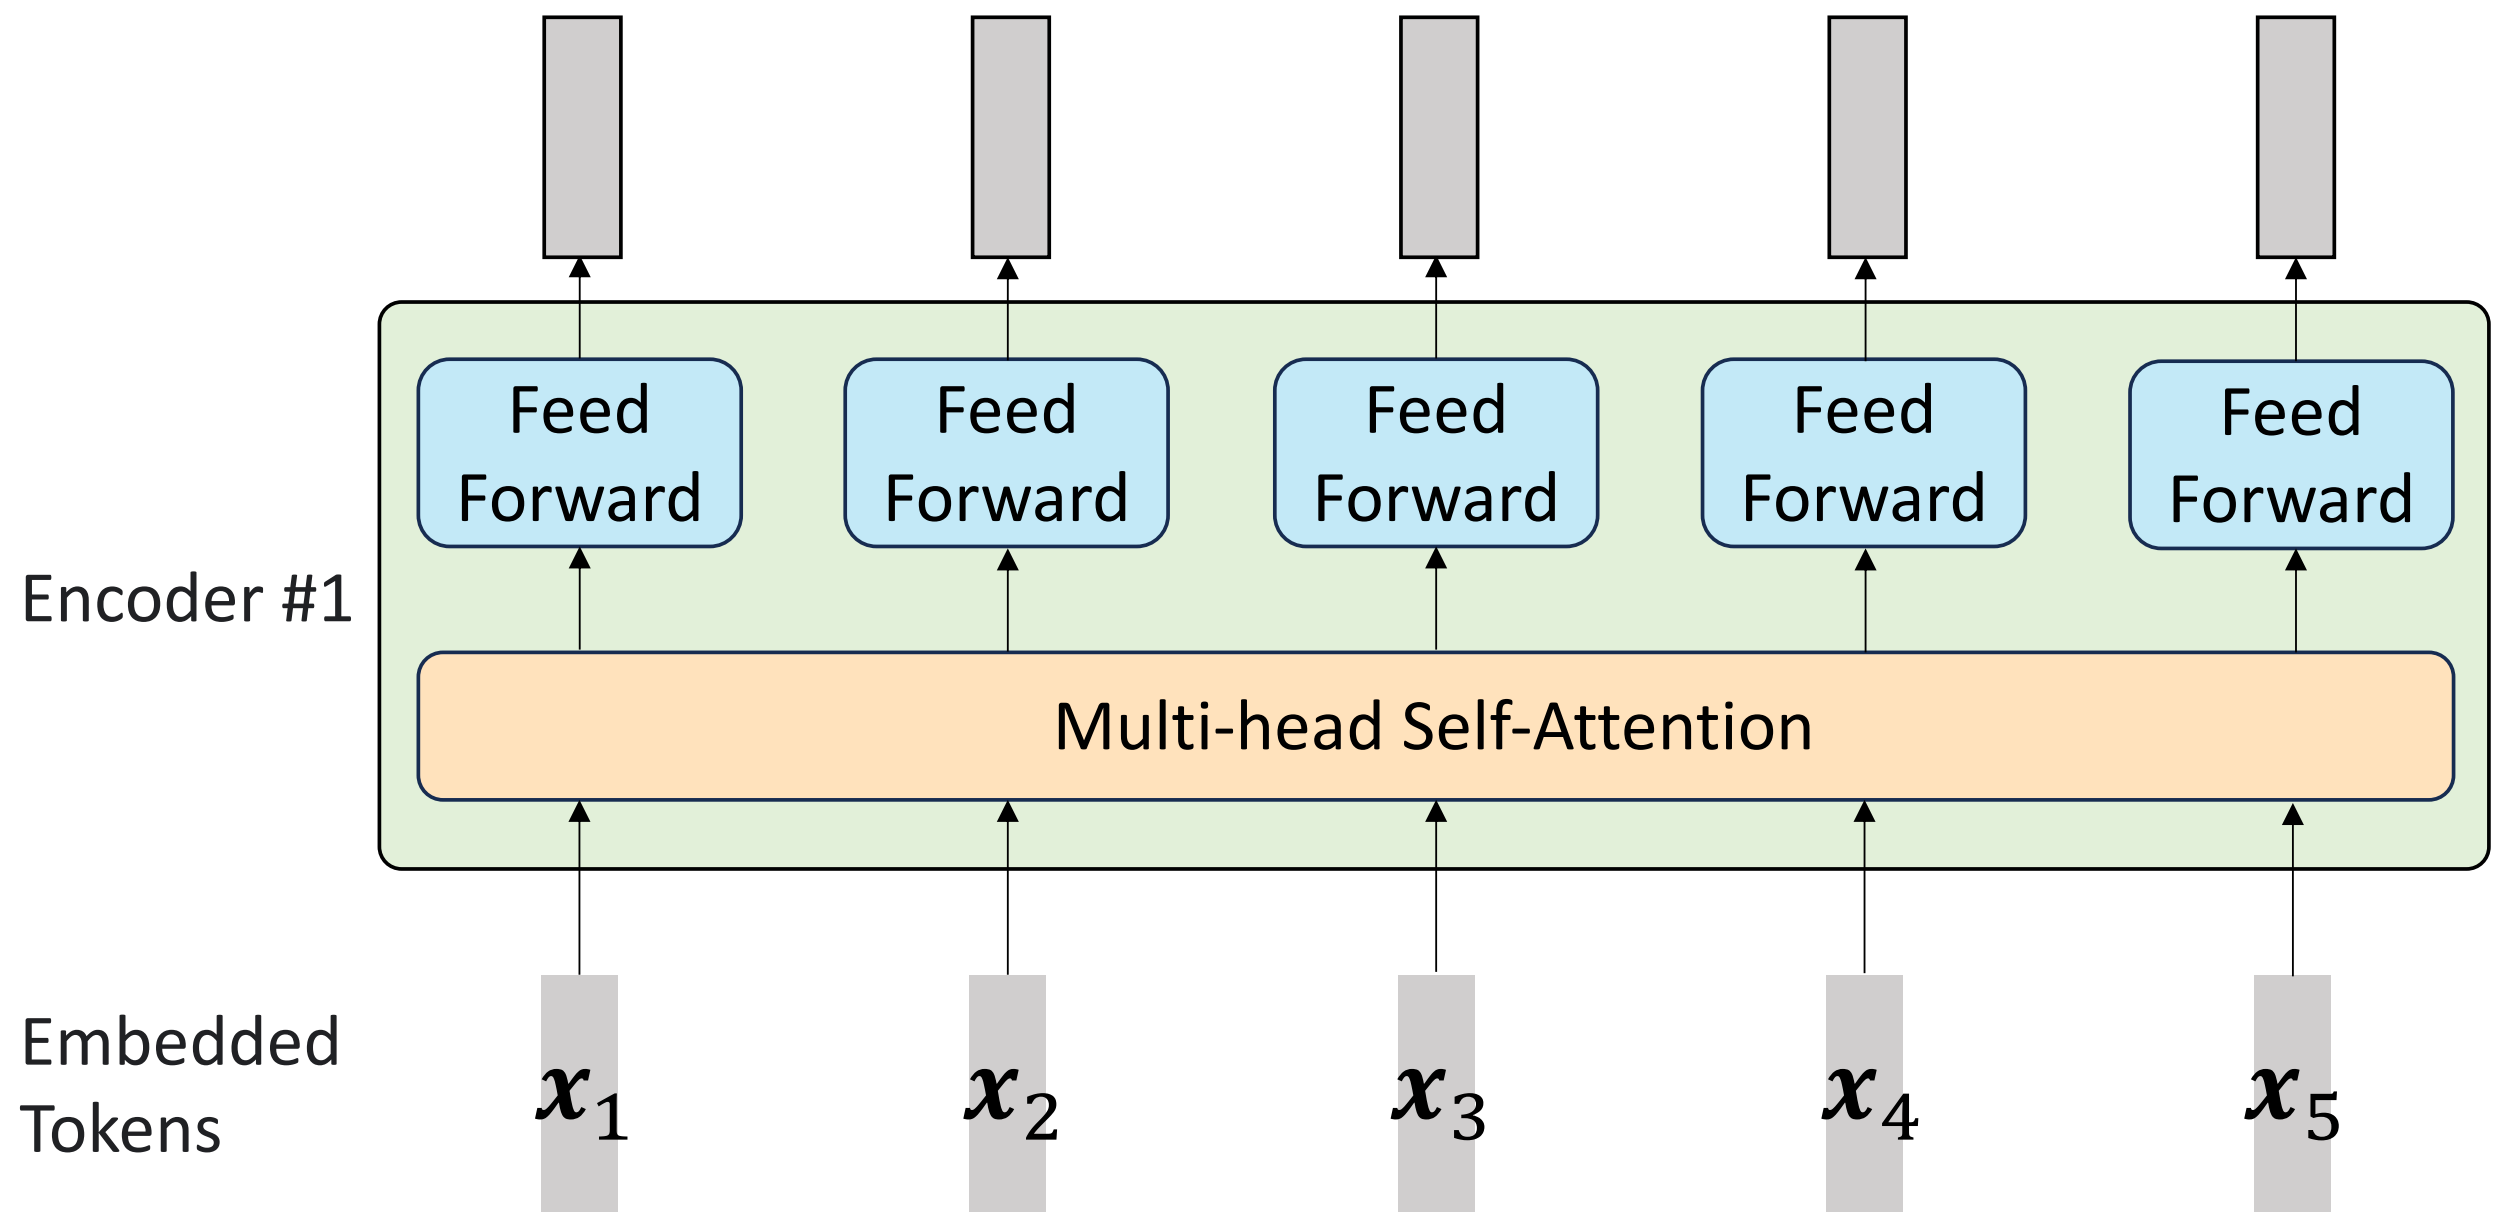
Self-Attention

Single-head attention

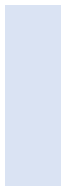
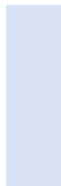
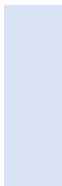
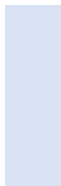
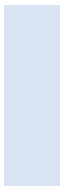





$$\text{Attention}(Q, K, V) = V \text{softmax}\left(\frac{K^\top Q}{\sqrt{d_k}}\right)$$

$$\begin{aligned} Q &= \boxed{W^Q} \begin{matrix} x_1 & x_2 & x_3 & x_4 & x_5 \end{matrix} \\ K &= \boxed{W^K} \begin{matrix} x_1 & x_2 & x_3 & x_4 & x_5 \end{matrix} \\ V &= \boxed{W^V} \begin{matrix} x_1 & x_2 & x_3 & x_4 & x_5 \end{matrix} \end{aligned}$$

MLPs





Positional Encoding

Position k		0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	
Dimension	2^3	0	0	0	0	0	0	0	0	1	1	1	1	1	1	1	1	← Slow oscillating
	2^2	0	0	0	0	1	1	1	1	0	0	0	0	1	1	1	1	
	2^1	0	0	1	1	0	0	1	1	0	0	1	1	0	0	1	1	
	2^0	0	1	0	1	0	1	0	1	0	1	0	1	0	1	0	1	← Fast oscillating
Positional embedding		$\updownarrow d$																
Embedded Tokens		$\updownarrow d$		 \mathbf{x}_1		 \mathbf{x}_2		 \mathbf{x}_3		 \mathbf{x}_4		 \mathbf{x}_5						

Positional Encoding

$$\mathbf{S}_i = [\sin(\omega_1 i), \cos(\omega_1 i), \sin(\omega_2 i), \cos(\omega_2 i), \dots, \sin(\omega_{e/2} i), \cos(\omega_{e/2} i)]$$



Normalized Range

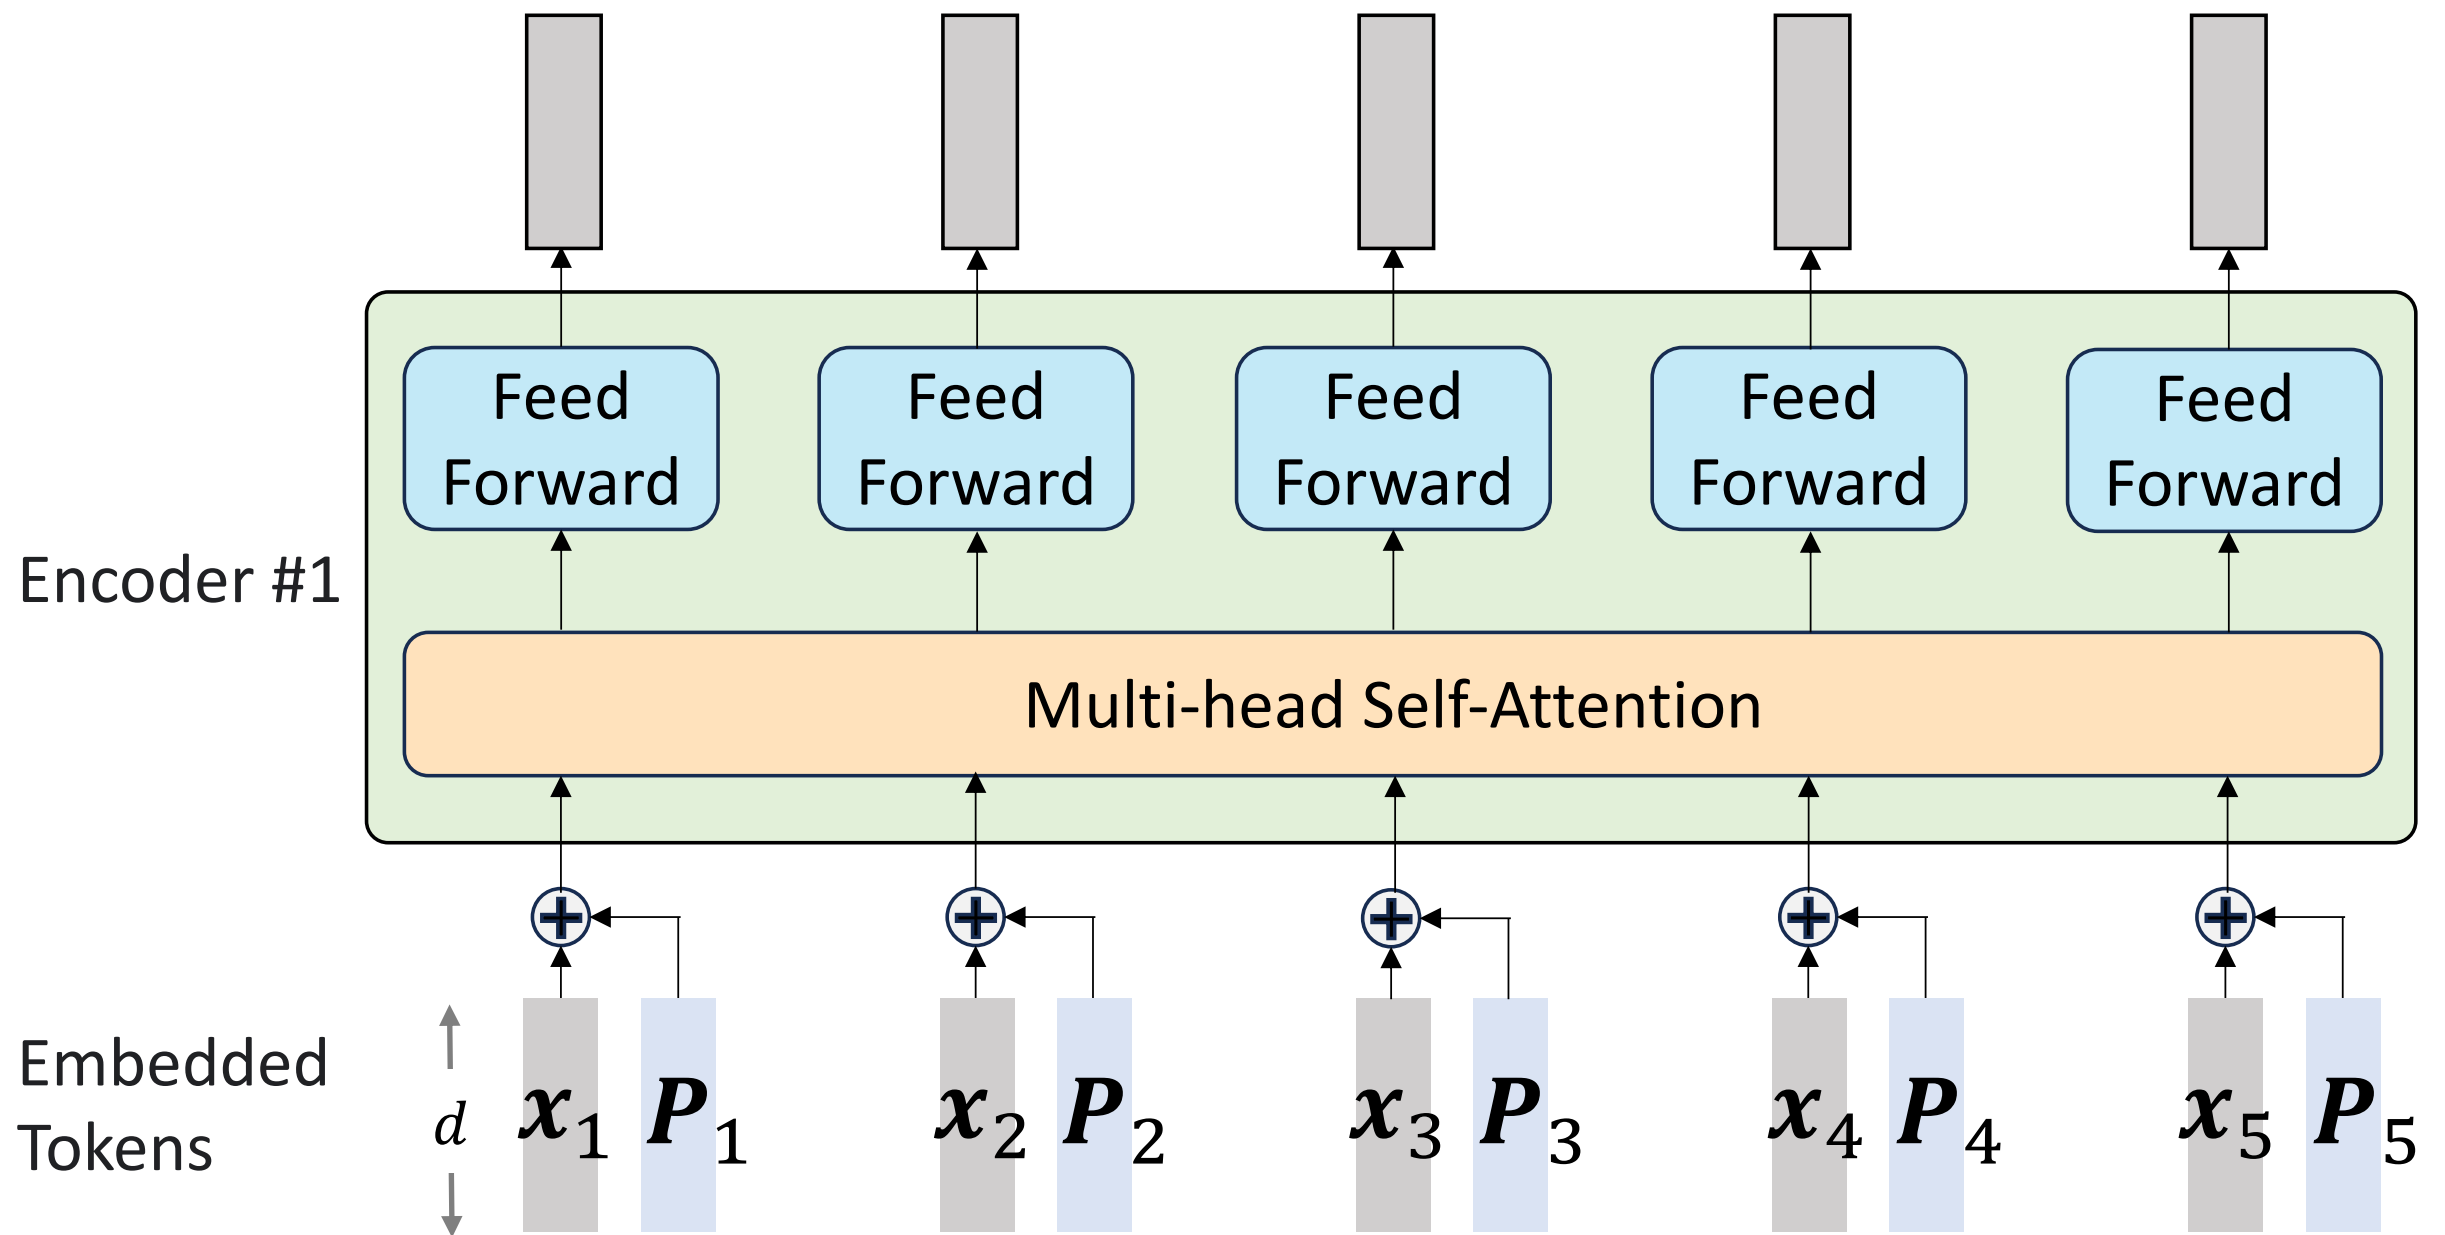


Unique identifier, unlimited length

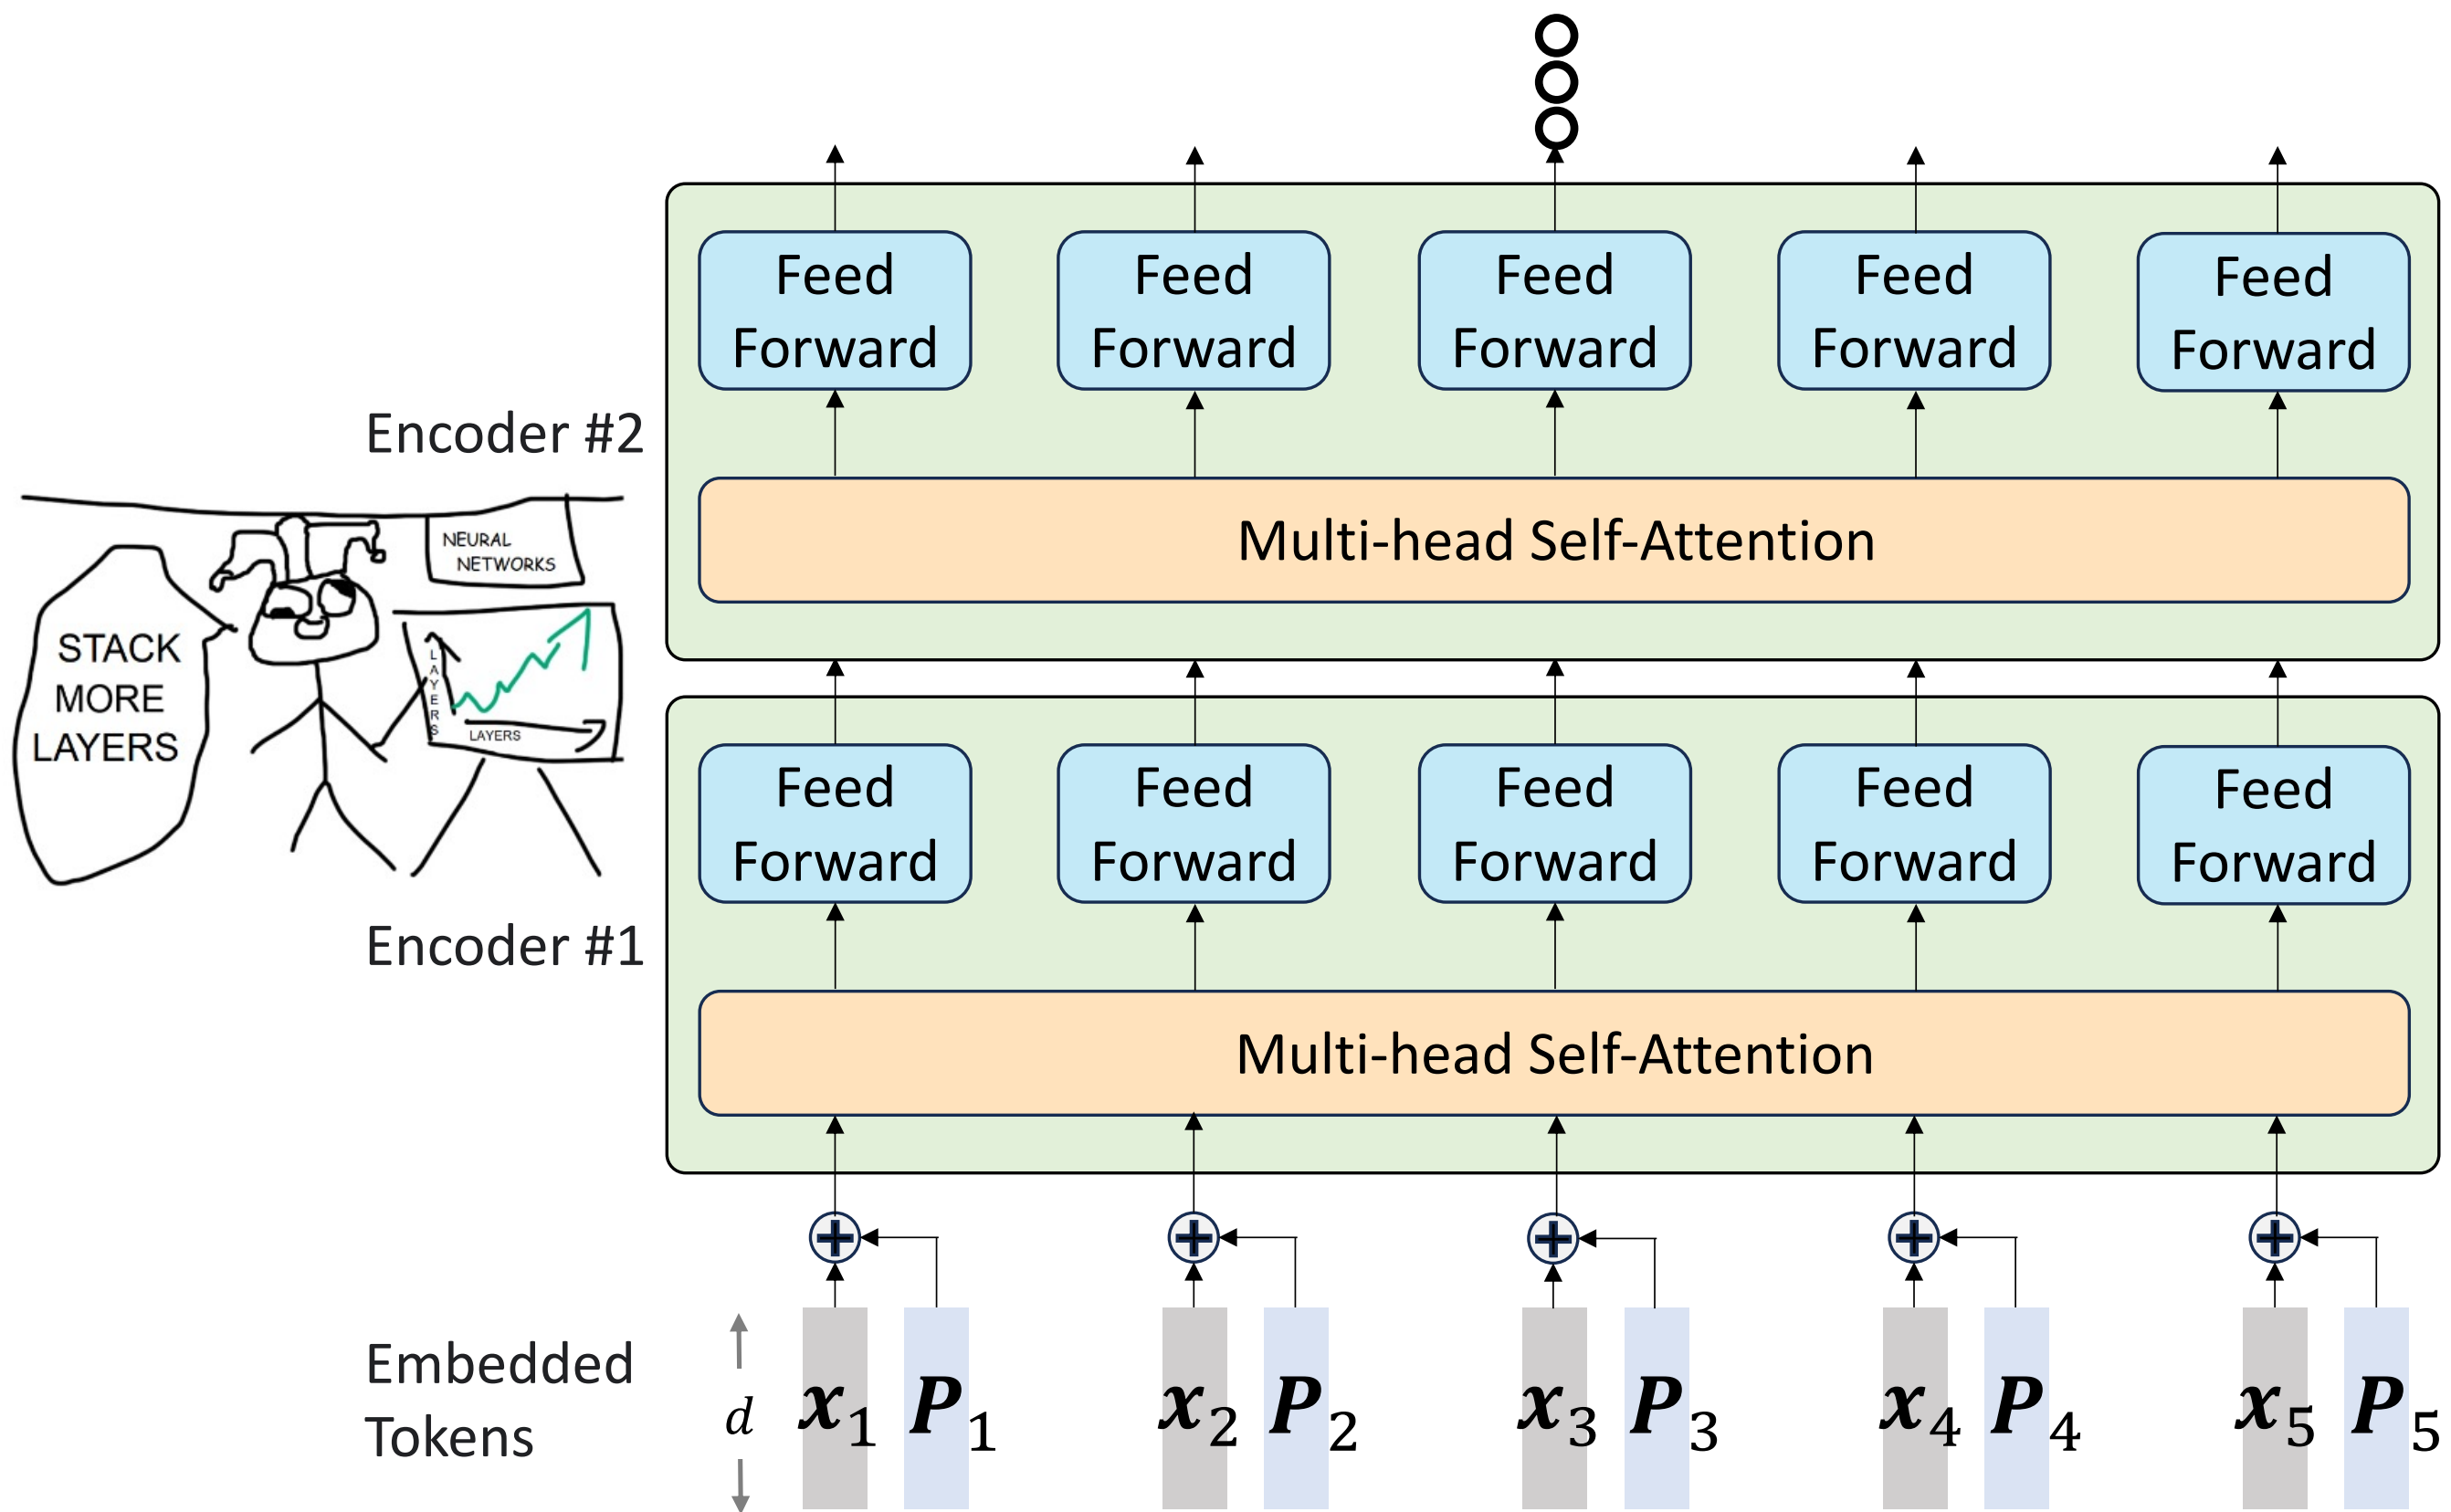


Relative positions as linear transform

Positional Encoding



Transformer Encoder



Transformer Encoder

💡 Residual connection

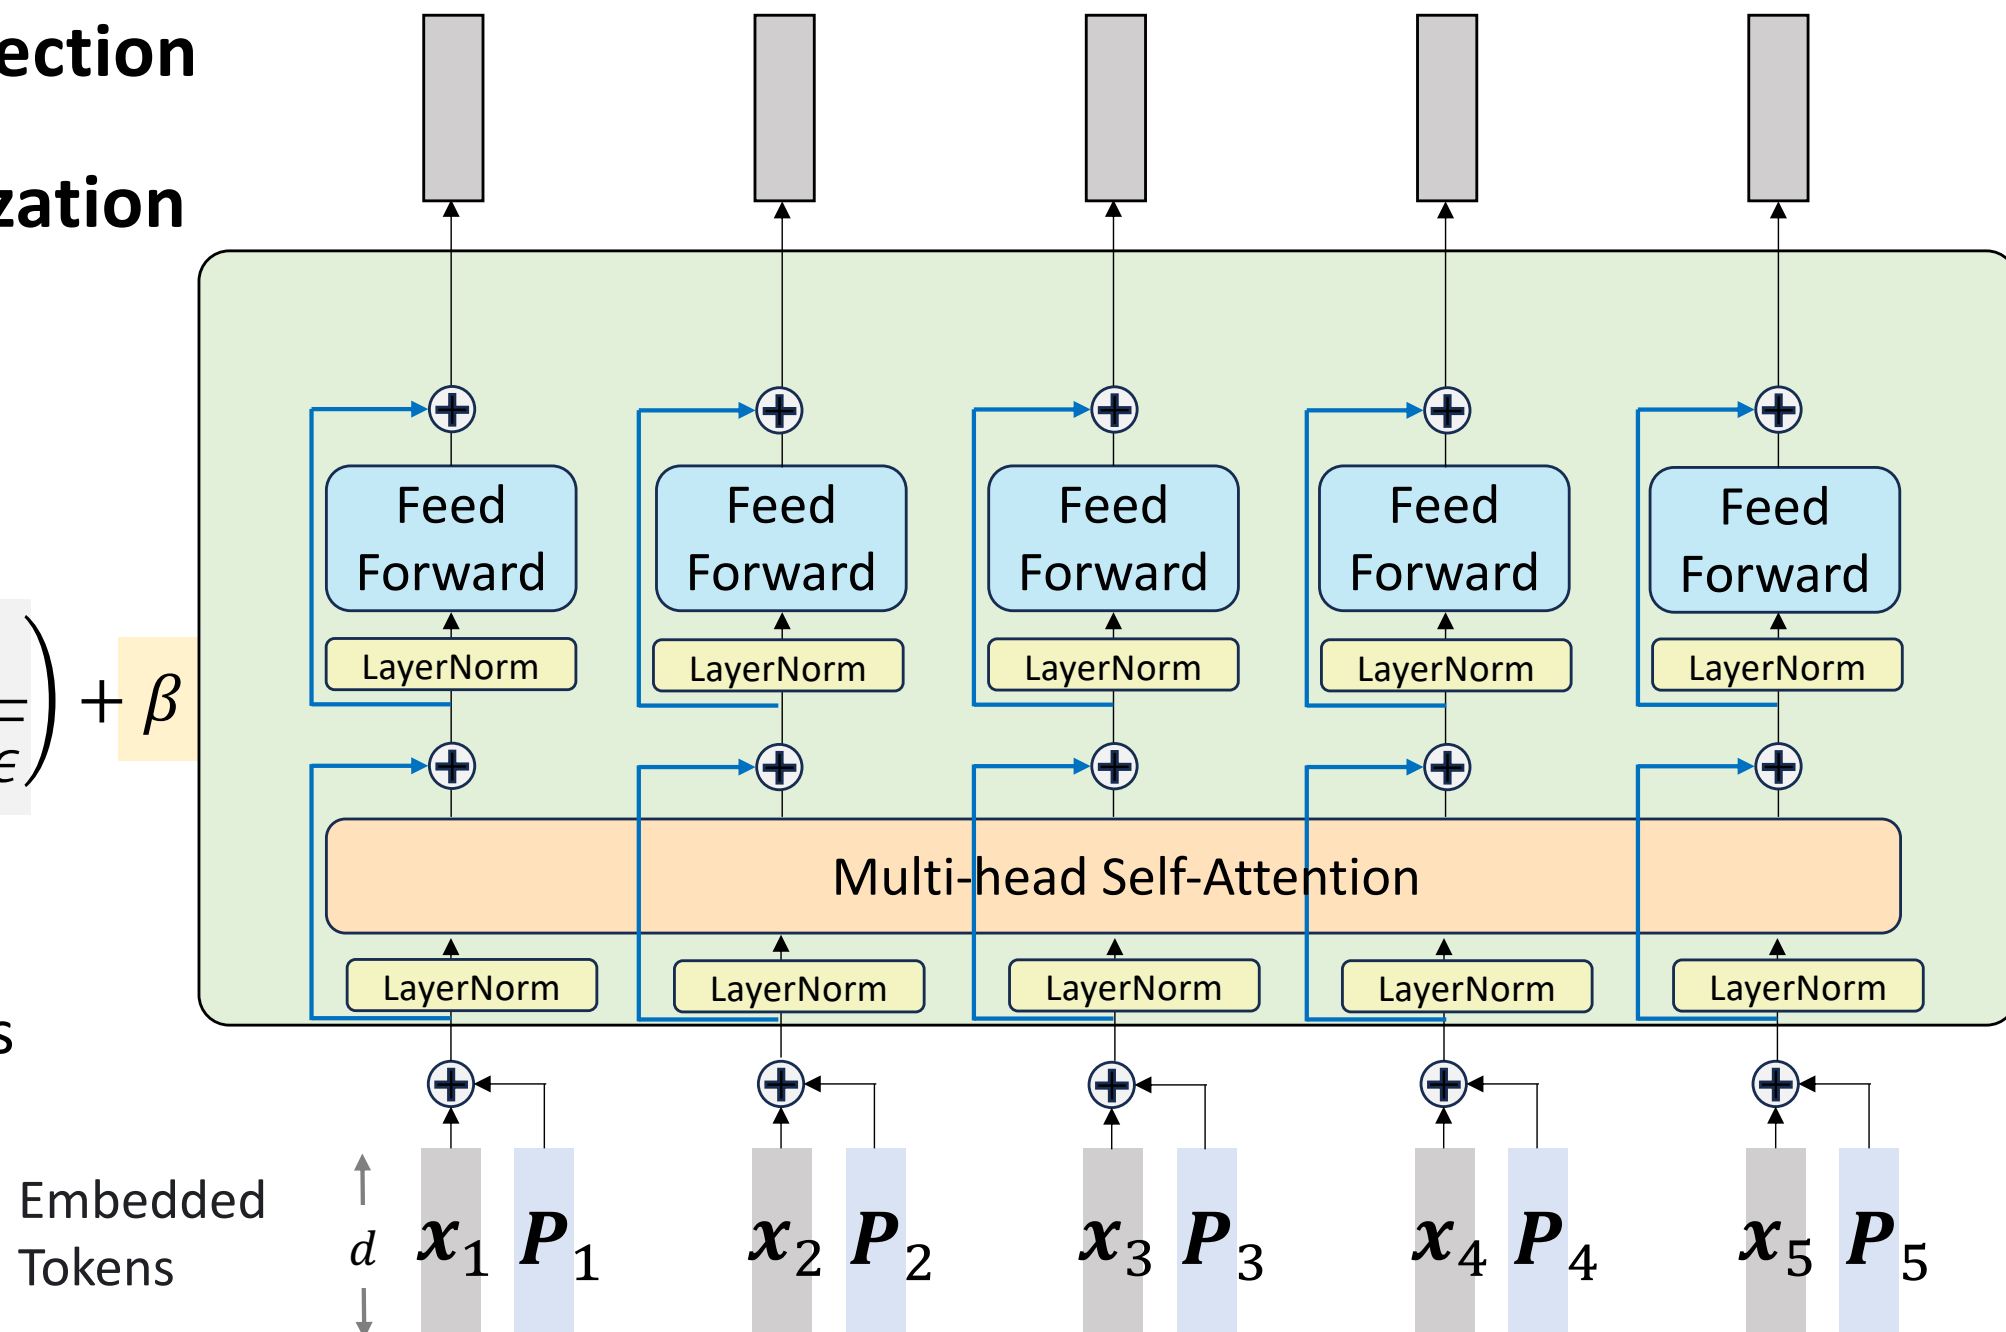
💡 Layer normalization

LayerNorm(\mathbf{x}) =

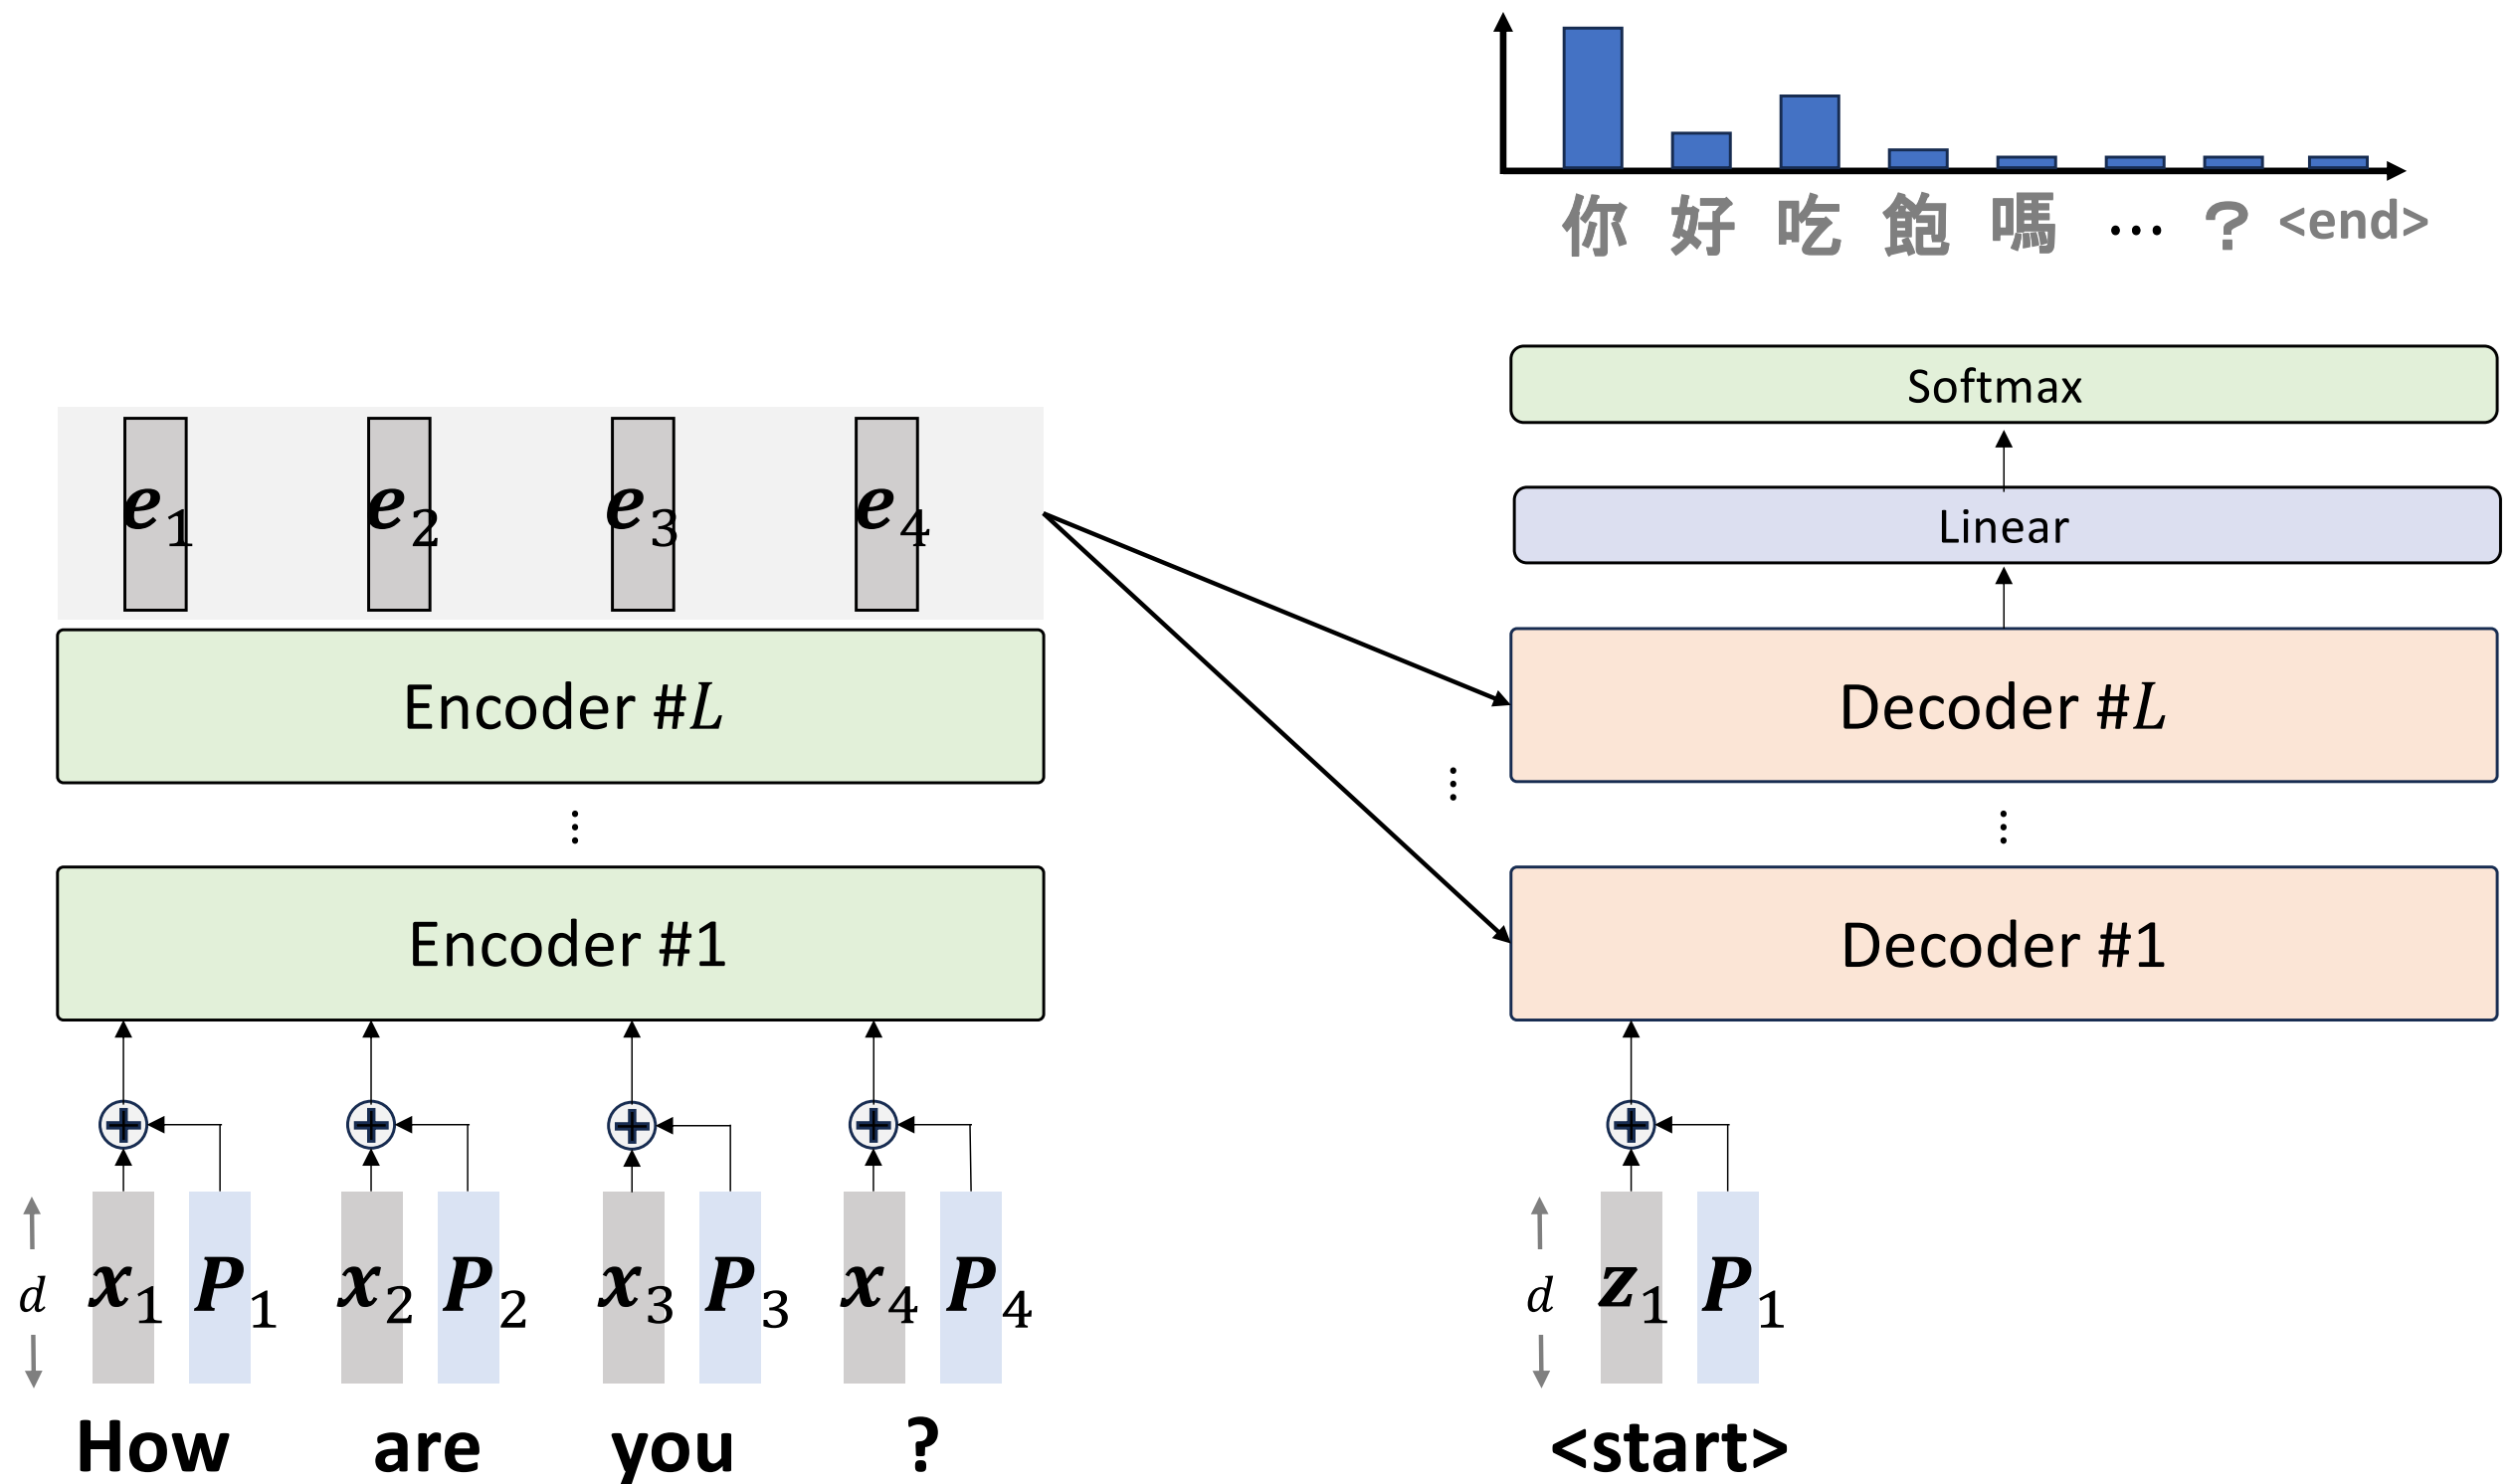
$$\gamma \left(\frac{\mathbf{x} - \text{mean}(\mathbf{x})}{\sqrt{\text{Variance}(\mathbf{x}) + \epsilon}} \right) + \beta$$

$\gamma, \beta \in R$

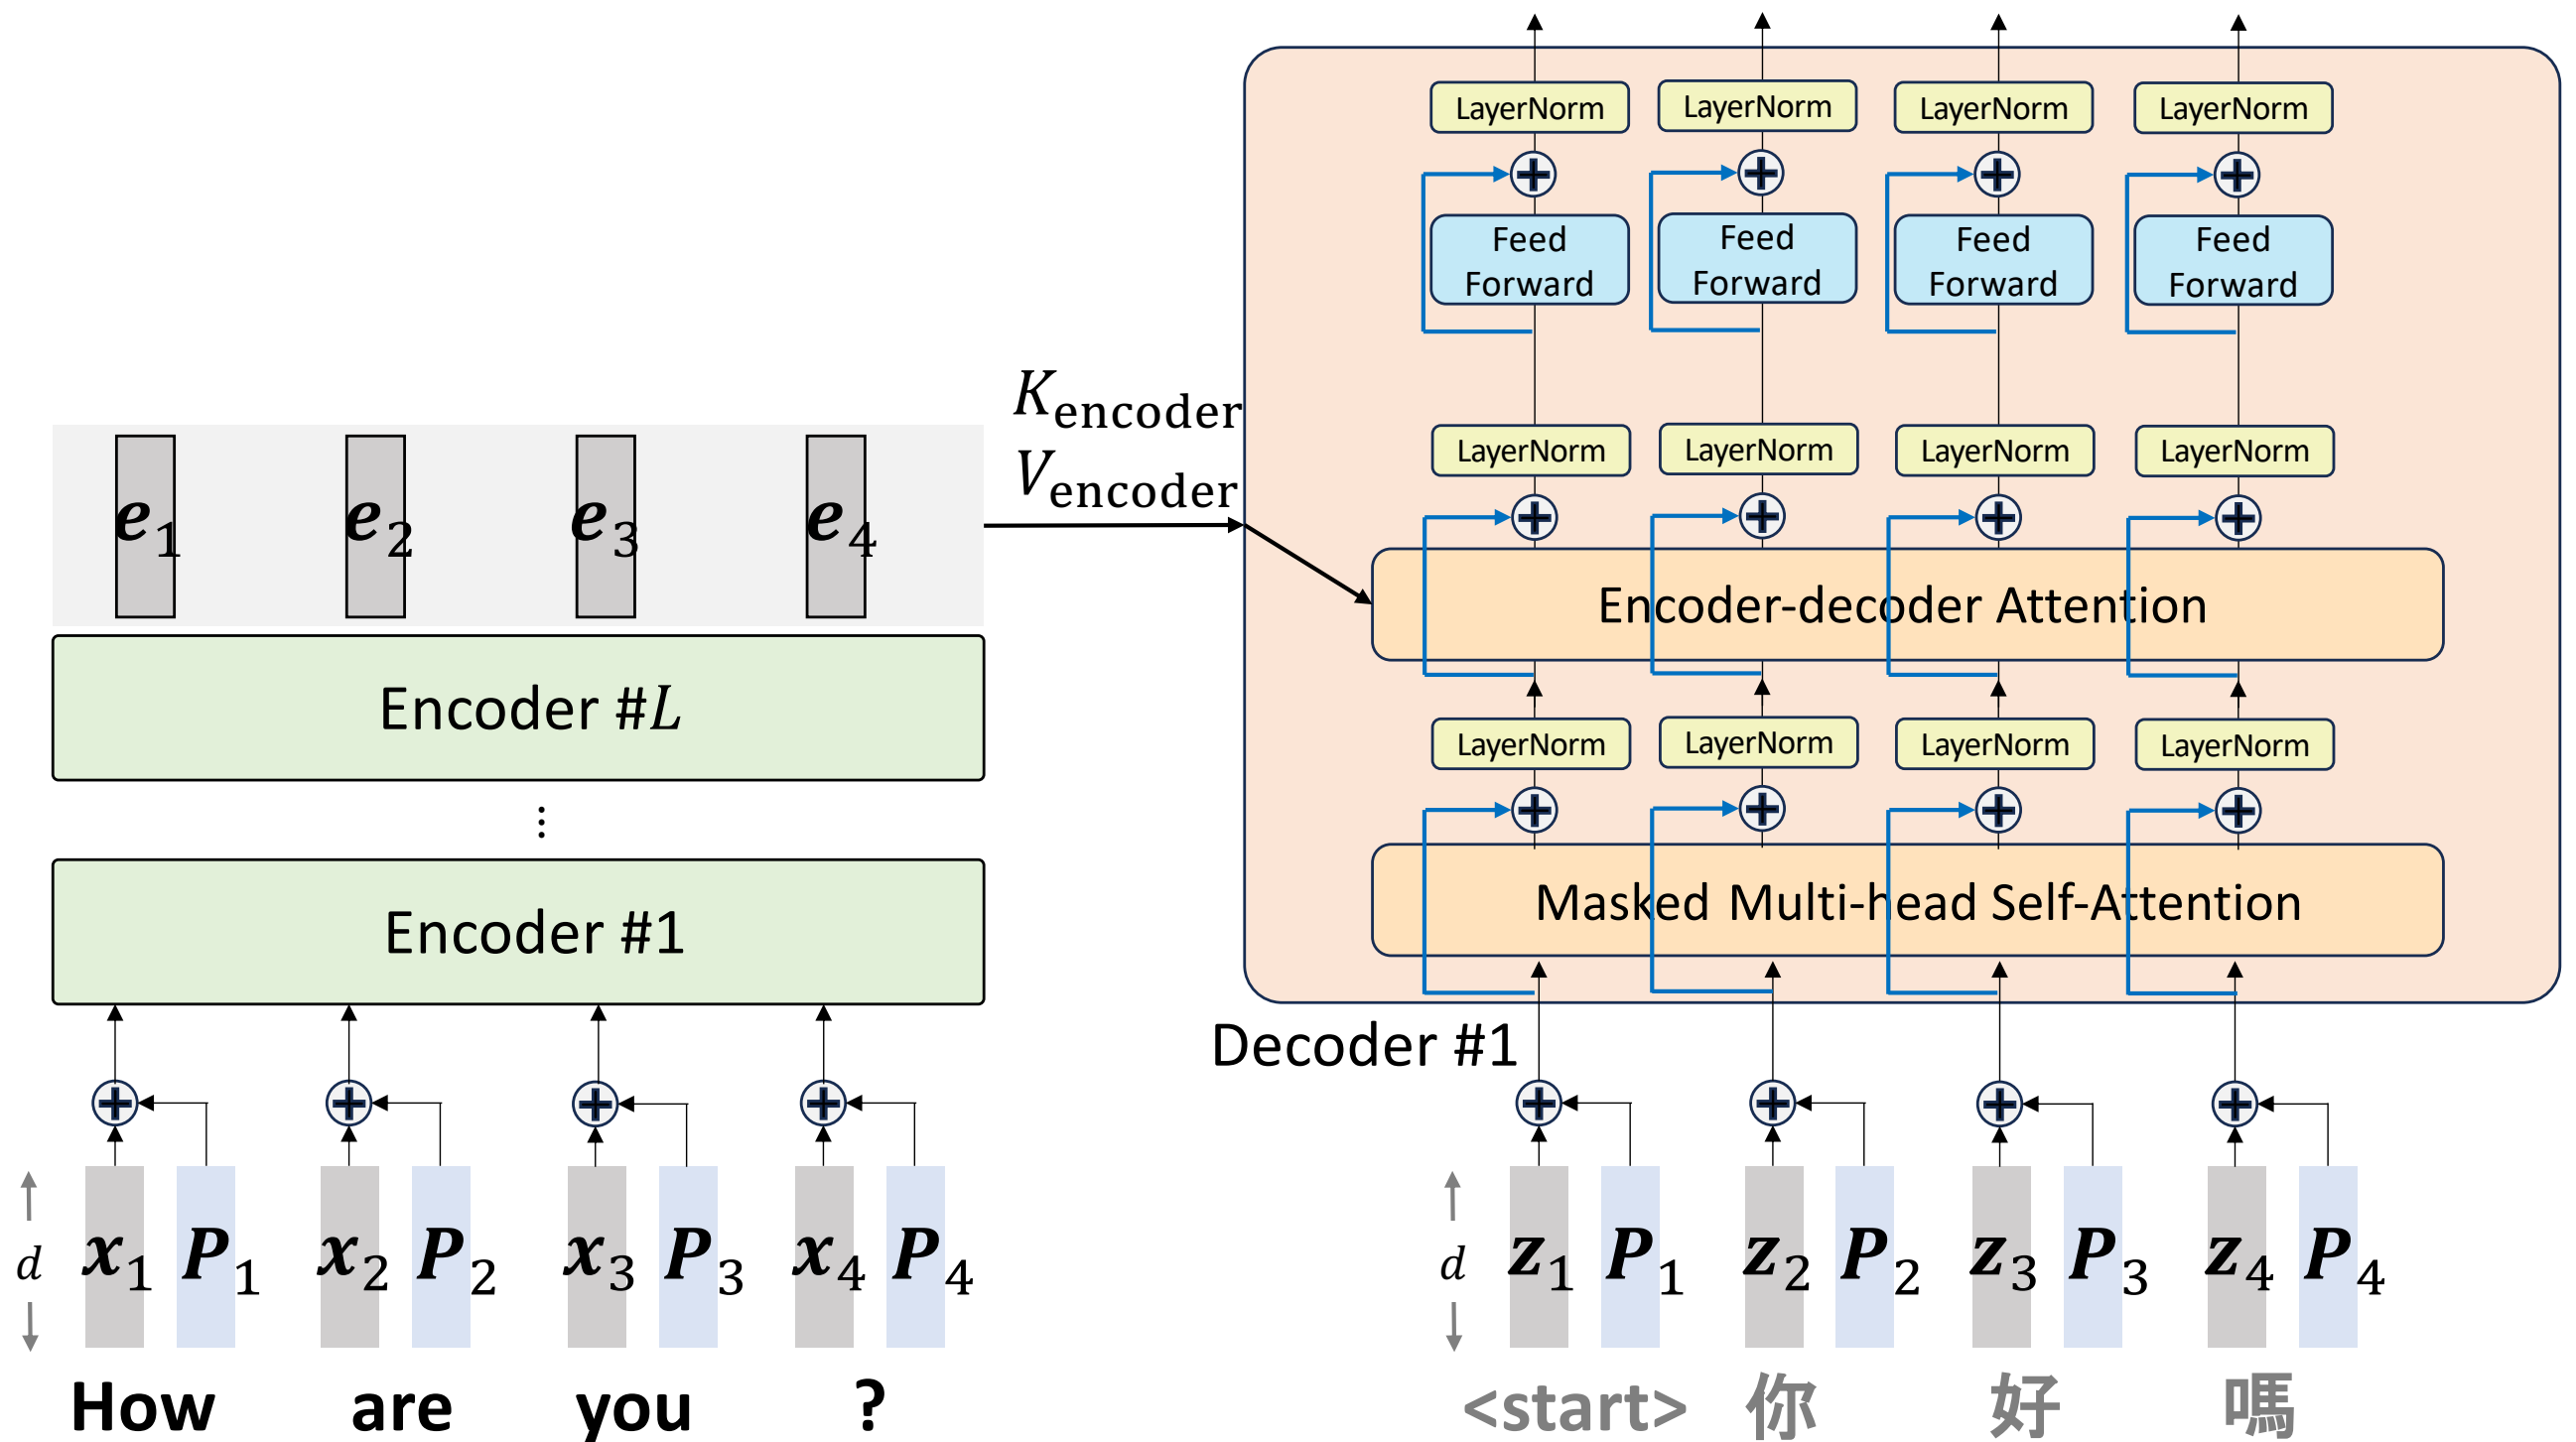
Learnable parameters



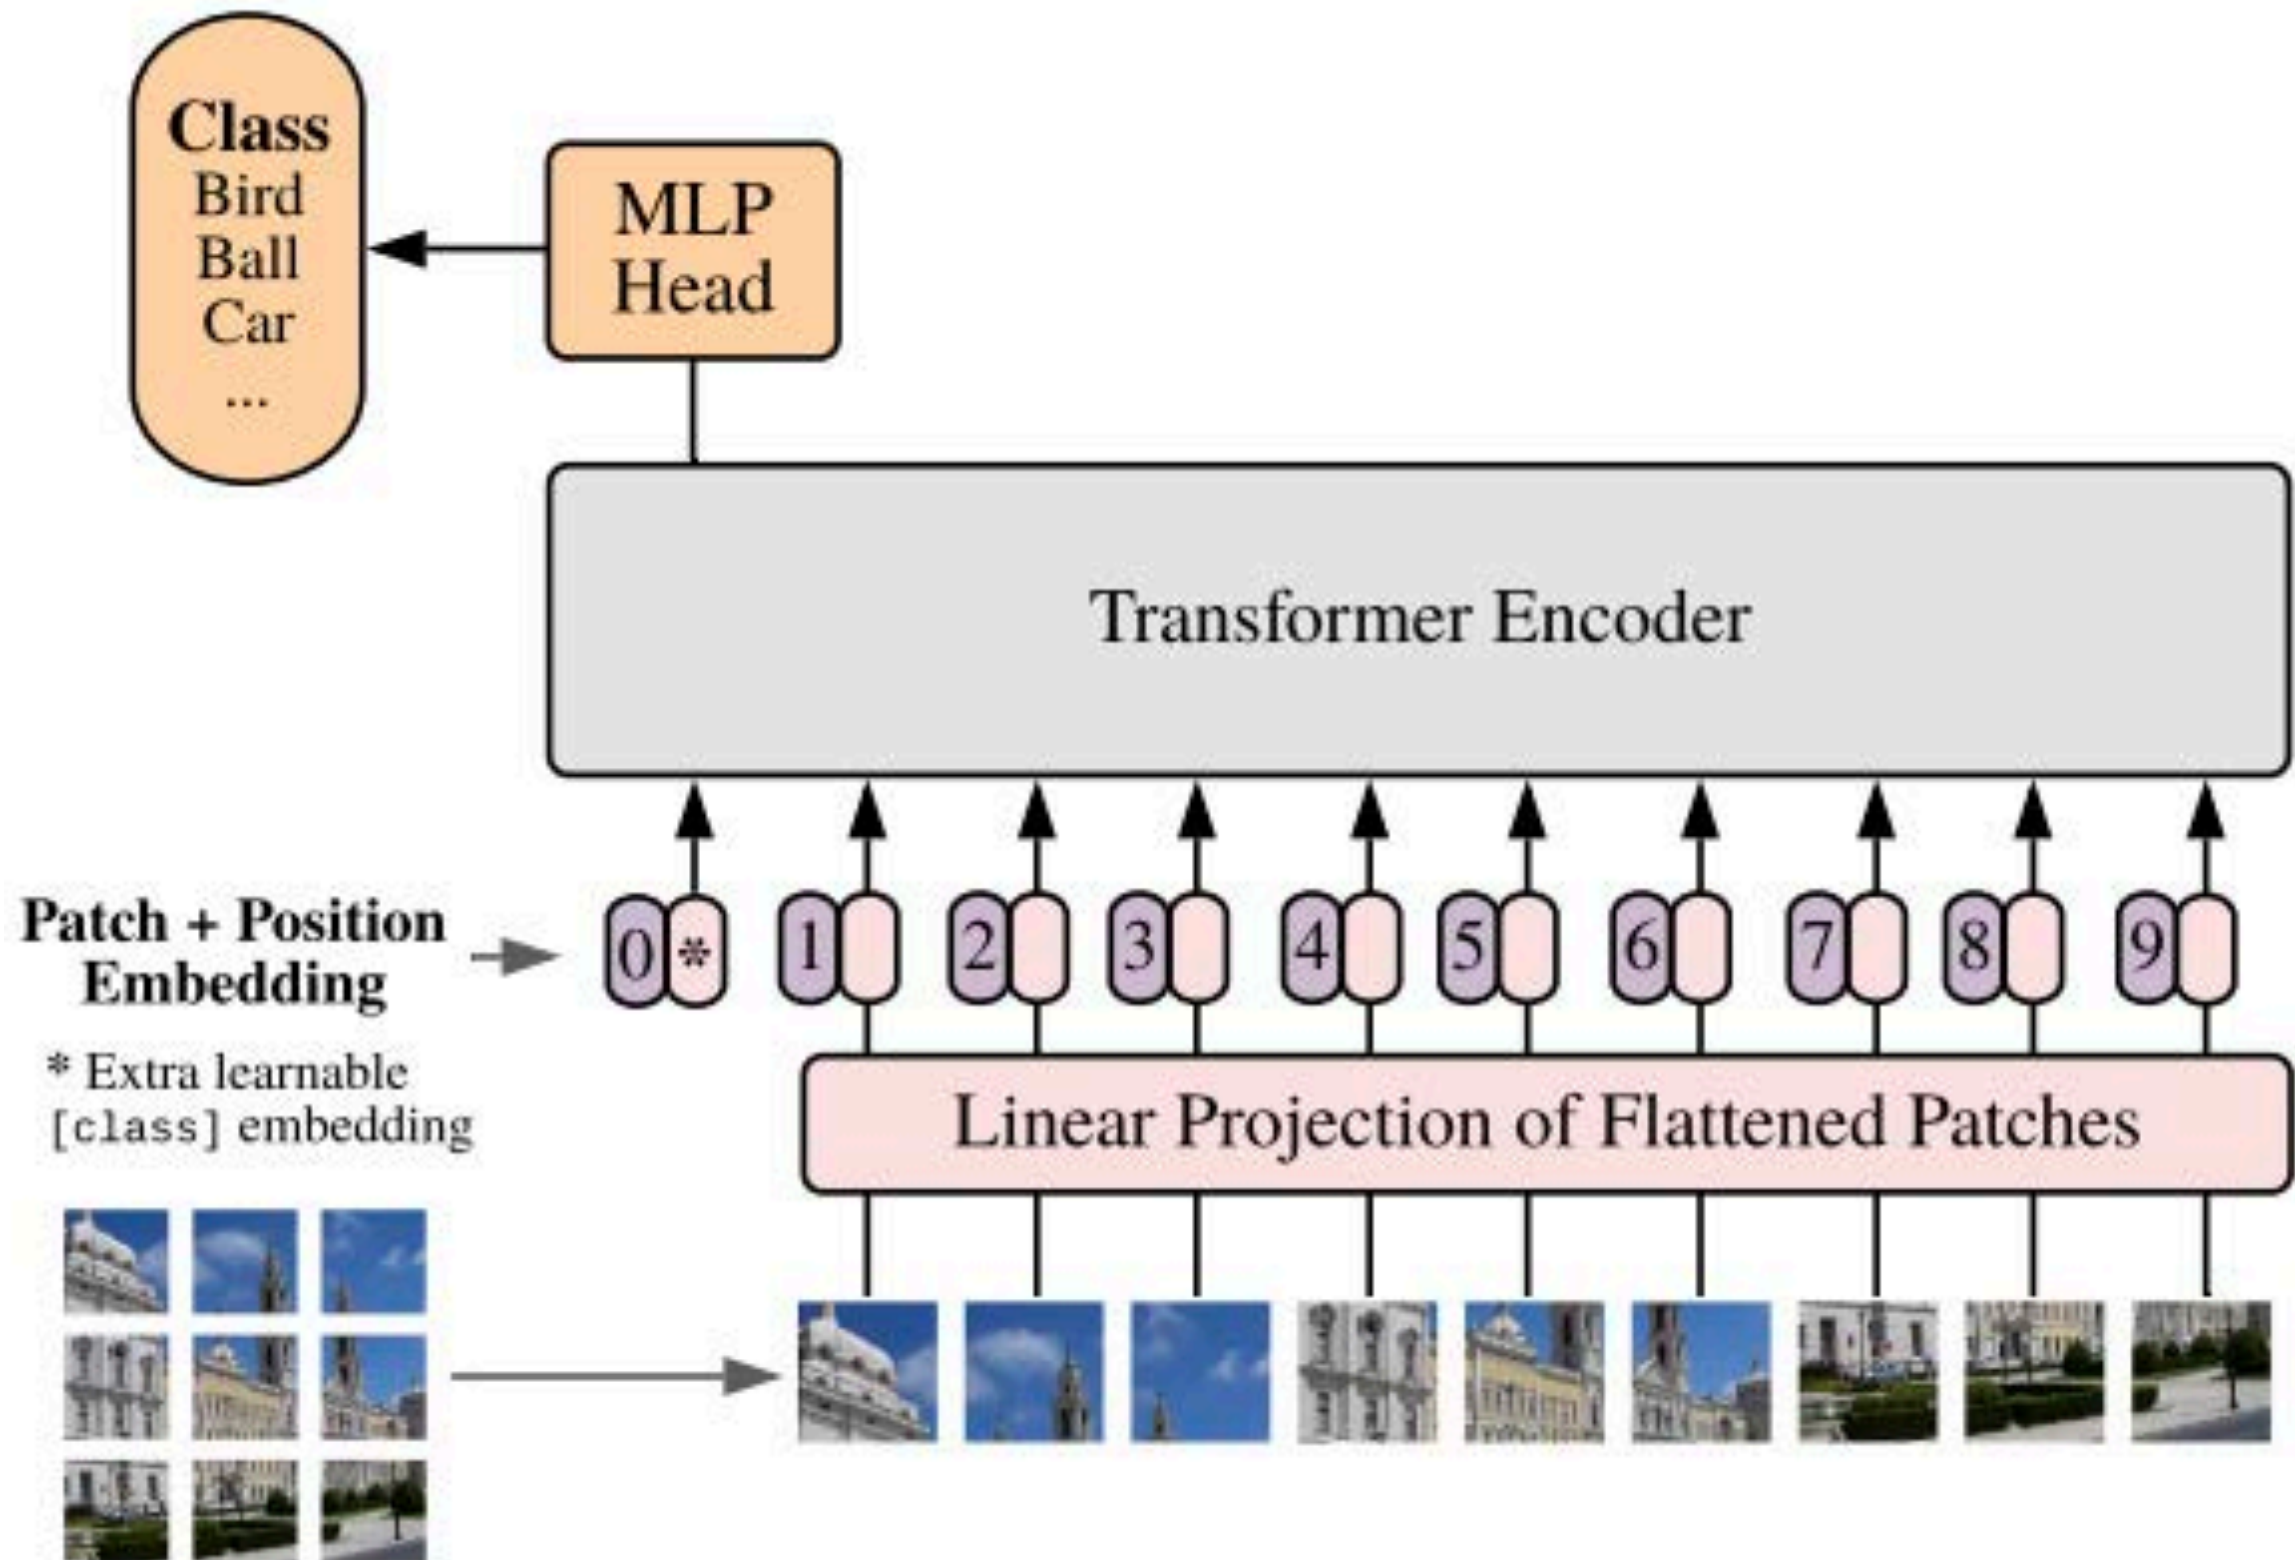
Decoders



Decoders

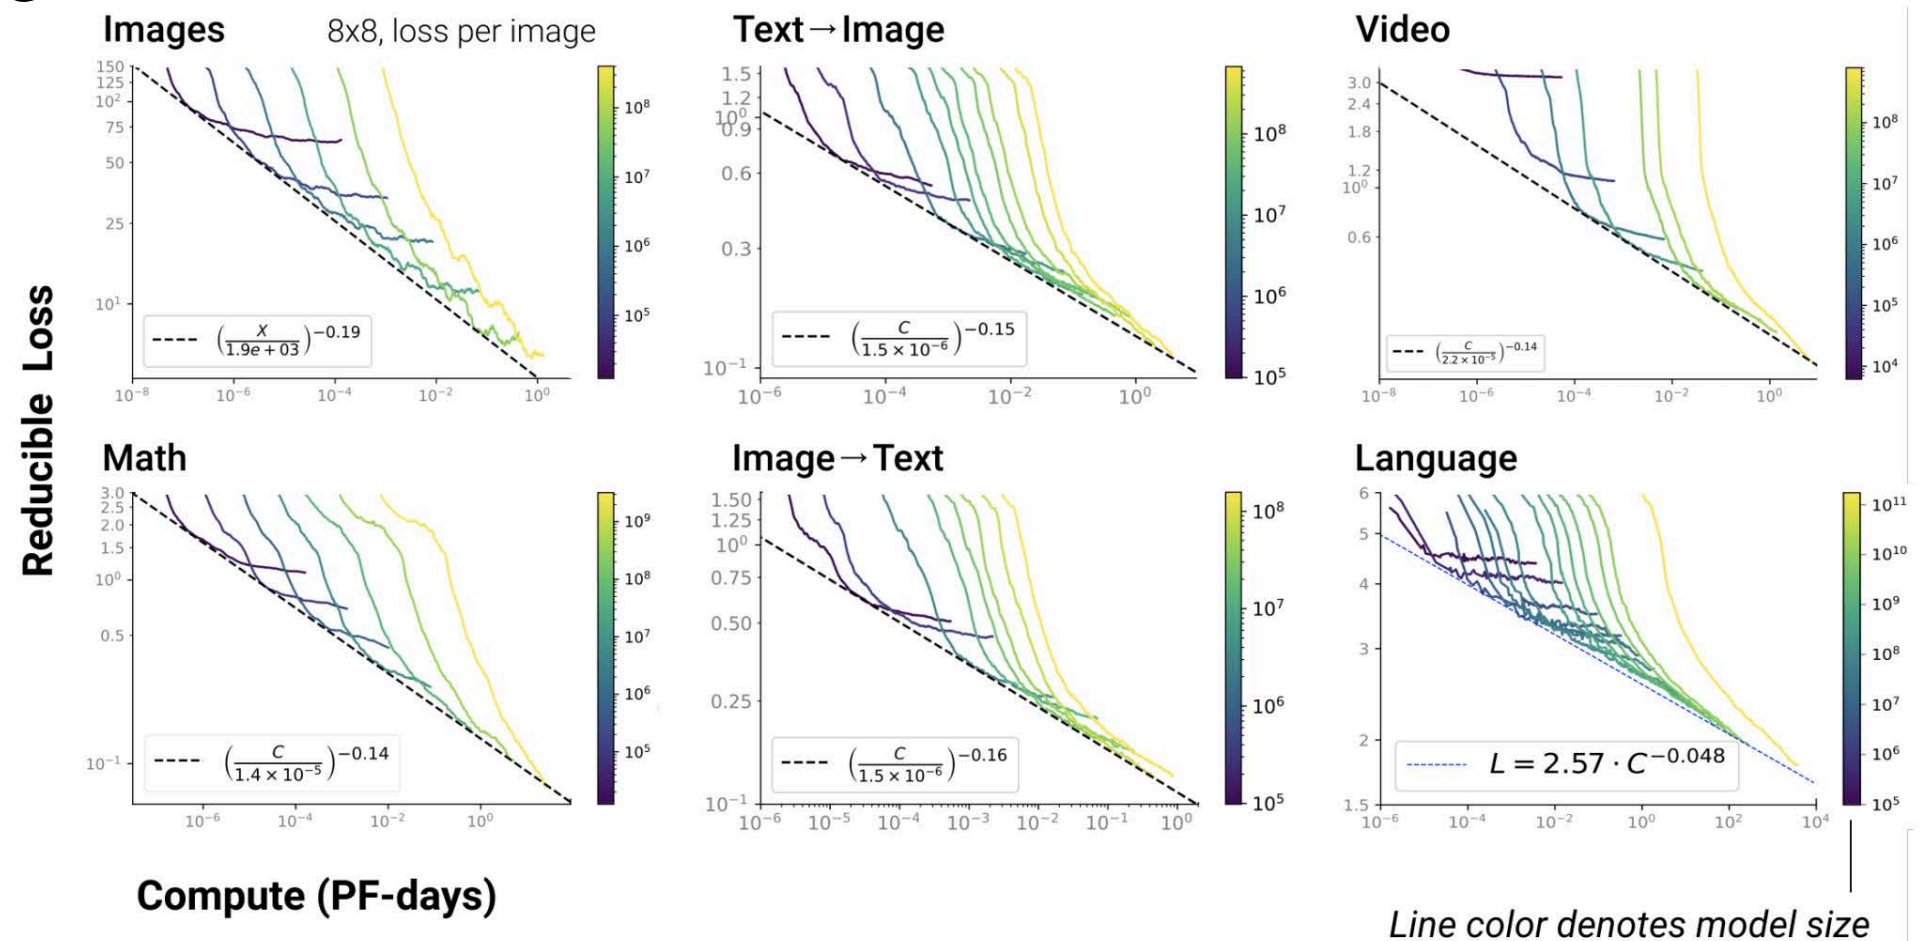


Vision Transformers



Advantages & Limitations

- Parallel processing
- Better at capturing long-range dependencies
- No vanishing gradient problem
- Scales well with data and compute
- Modality-agnostic



- Self-attention cost and memory scales quadratically with seq. length
- Needs a lot of data
- Expensive and difficult to train

Generative Pretrained Transformers

New chat

chat.openai.com/chat

VPN Archvizu Hexlet РАБОТА Competitions StackEdit c# Rhino Python Материал для ви... Полезные ссылки Network Free Content» Другие закладки

+ New chat

New chat

Blender Hello Popup Script

Clear conversations

Dark mode

OpenAI Discord

Updates & FAQ

Log out

EL

blender python code which will create cube 10,10,10

Here is a code to create a cube with dimensions 10x10x10 in Python:

ruby

Copy code

```
class Cube:
    def __init__(self, side):
        self.side = side

    def volume(self):
        return self.side**3

cube = Cube(10)
print(cube.volume())
```

This code defines a `Cube` class with a constructor method `__init__` that takes `side` as an argument and sets it as an instance variable. The `volume` method calculates the cube's volume by cubing the `side` attribute. Finally, an instance of the `Cube` class is created with `side=10` and its volume is printed.


Regenerate response

ChatGPT Jan 9 Version. Free Research Preview. Our goal is to make AI systems more natural and safe to interact with. Your feedback will help us improve.

Multi-modal Transformers

VisualQnA

Example



Upload Images



Click to upload or drag and drop

SVG, PNG, JPG

or

Import from URL



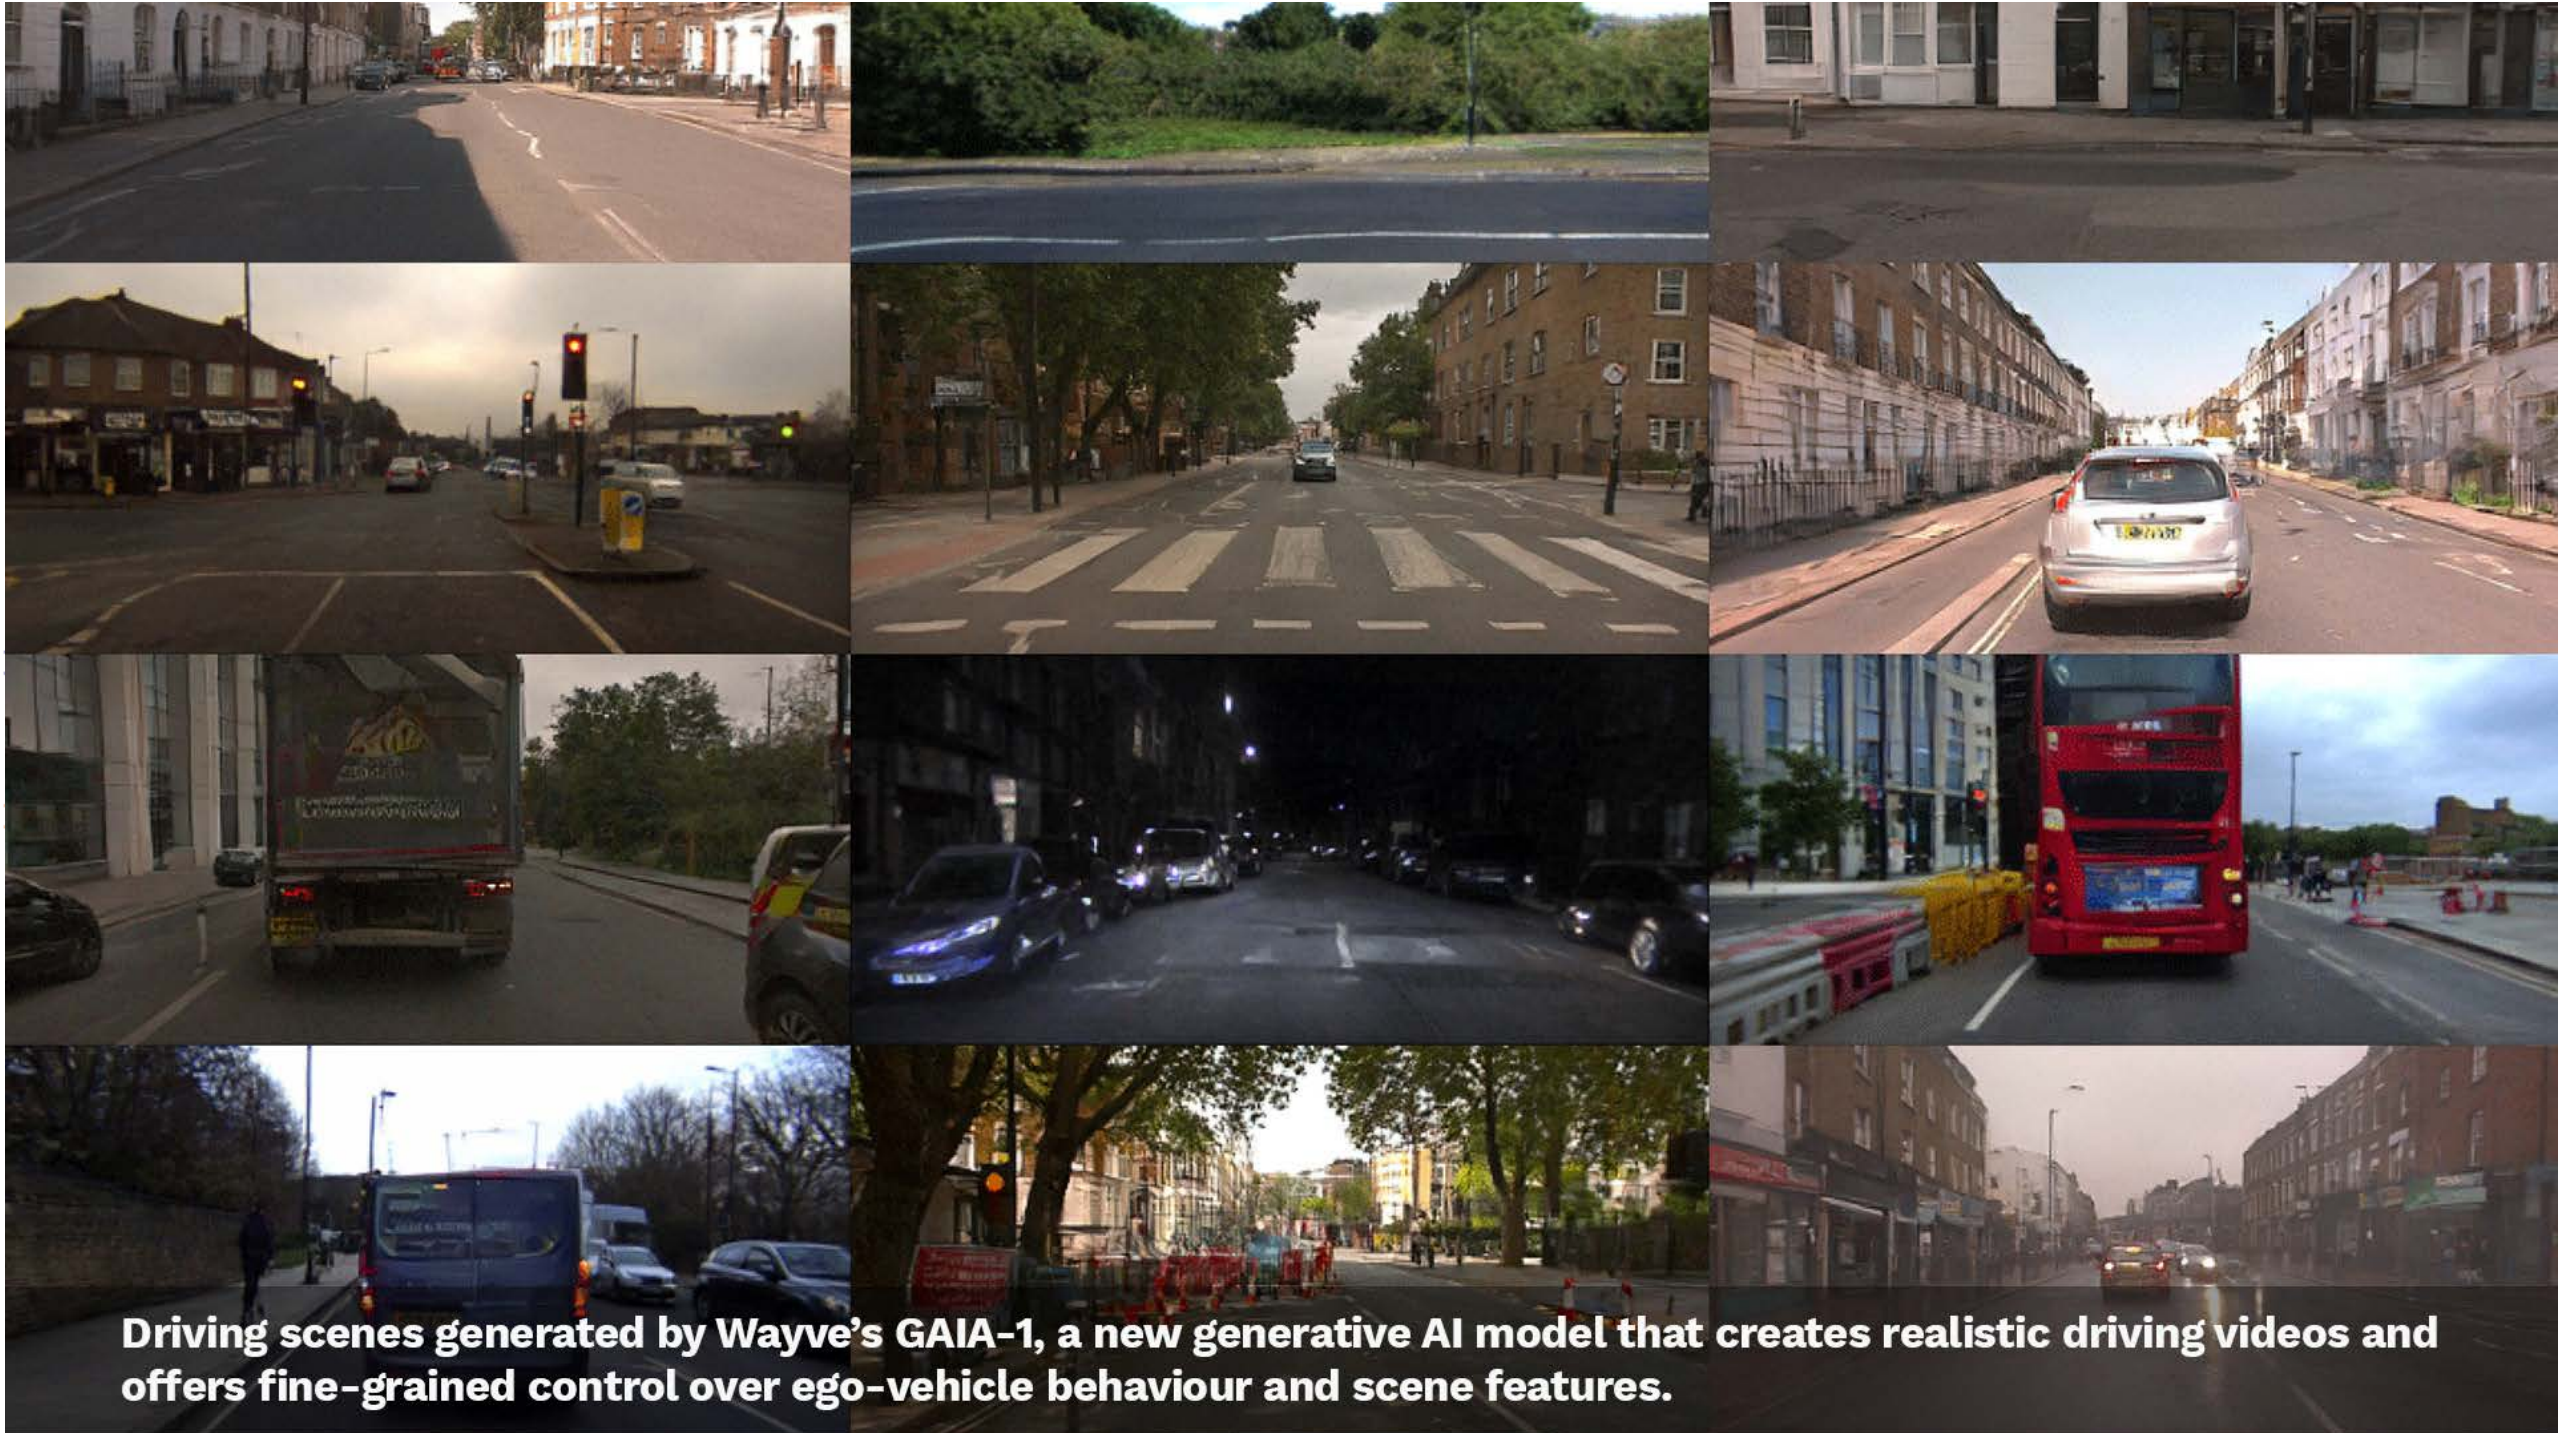
Parameters

Max output tokens: 128

Enter prompt here



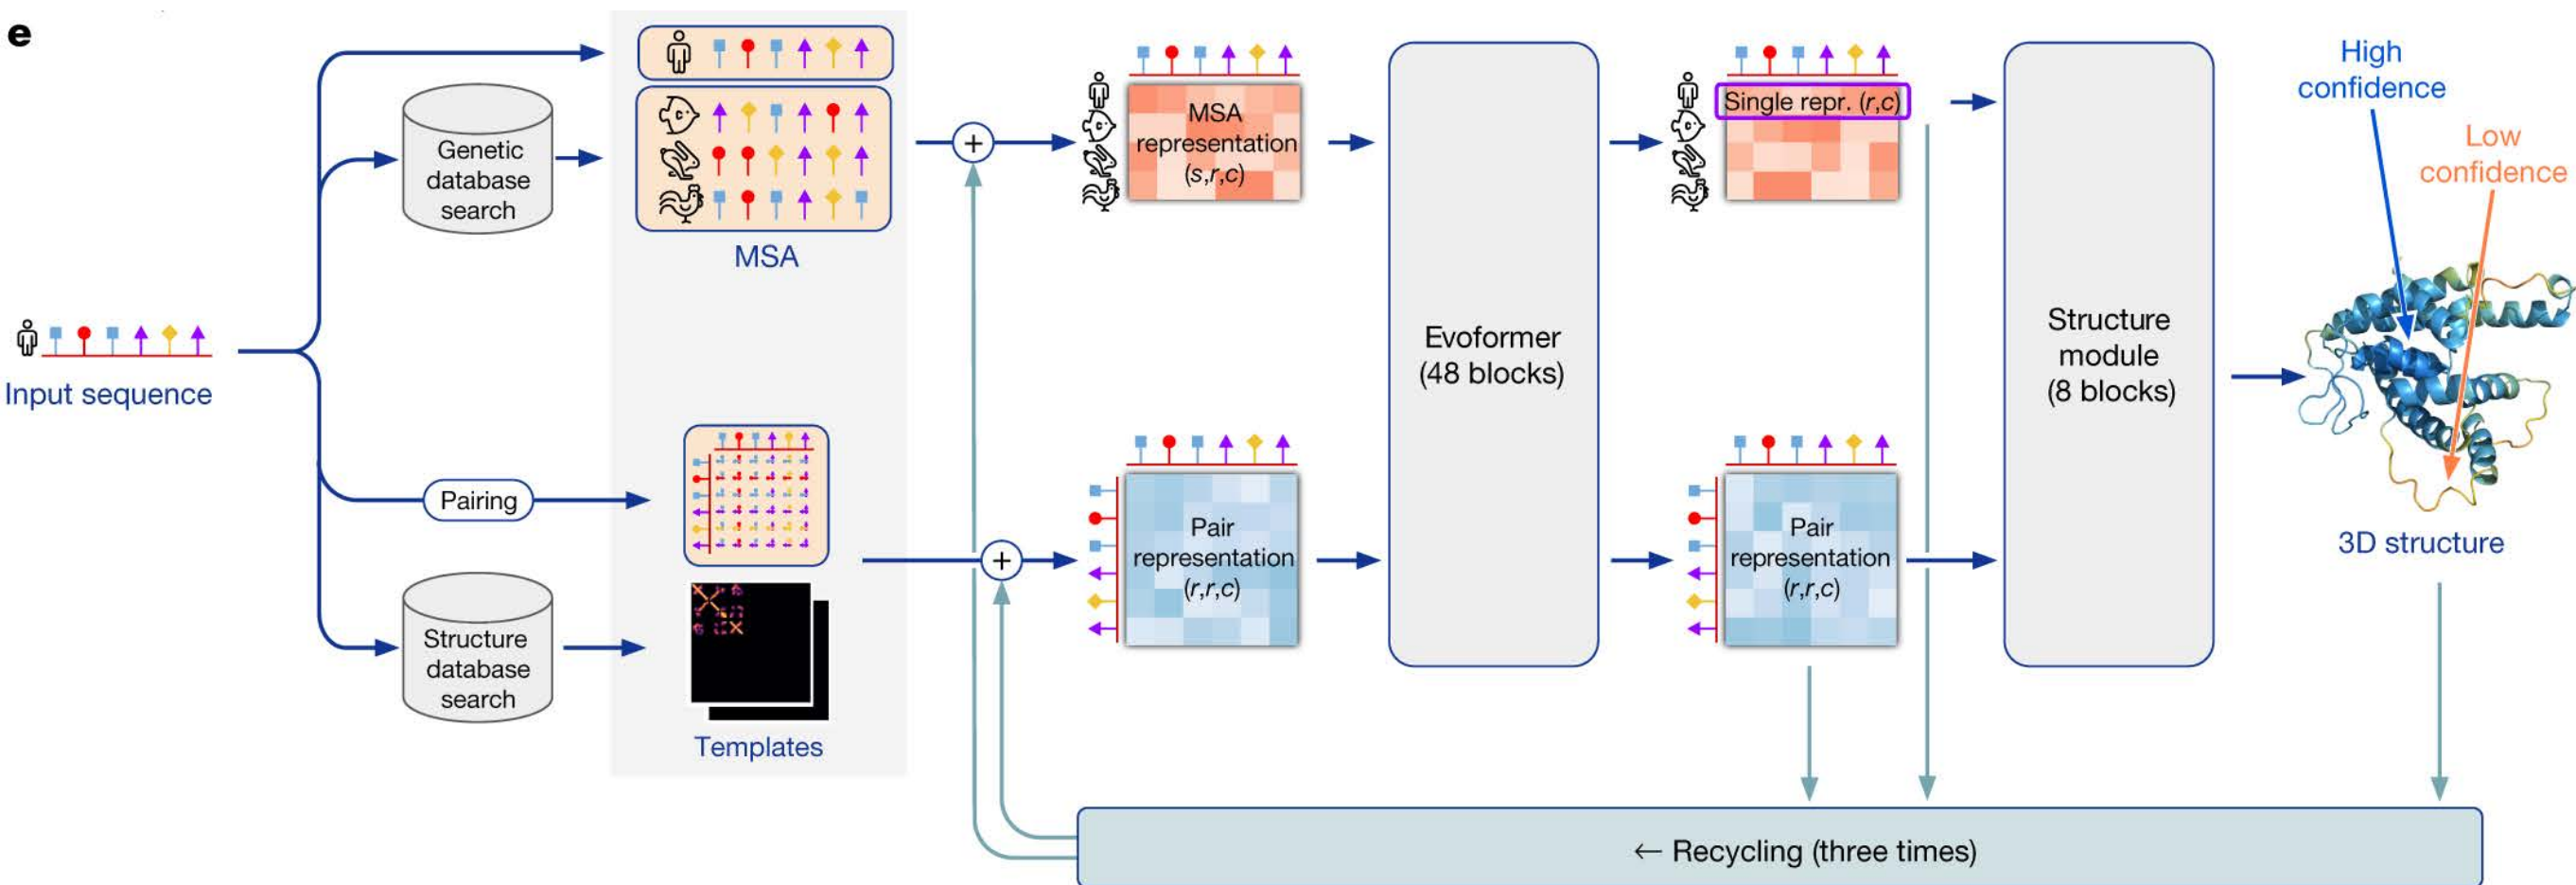
Multi-modal Transformers




Driving scenes generated by Wayve's GAIA-1, a new generative AI model that creates realistic driving videos and offers fine-grained control over ego-vehicle behaviour and scene features.

AlphaFold 3


e



7BBV



Ground truth shown in grey



Abramson, J., Adler, J., Dunger, J., Evans, R., Green, T., Pritzel, A., ... & Jumper, J. M. (2024). Accurate structure prediction of biomolecular interactions with AlphaFold 3. *Nature*, 1-3.

UvA Tutorial Notebooks

Build your own GPT from scratch

**LET'S BUILD GPT.
FROM SCRATCH.
IN CODE.
SPELLED OUT.**

